



3D Point Clouds

Lecture 2

Nearest Neighbors

主讲人 黎嘉信

Aptiv 自动驾驶
新加坡国立大学 博士
清华大学 本科





1. Binary Search Tree



2. Kd-tree



3. Octree

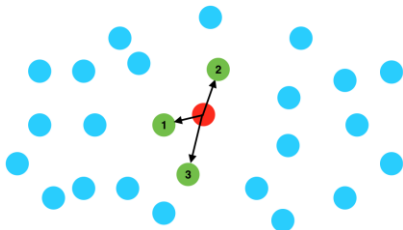


Nearest Neighbor (NN) Problem



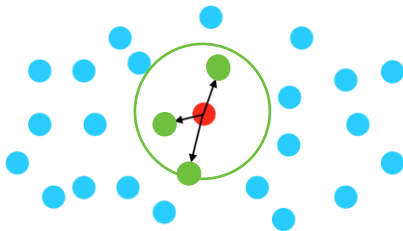
K-NN

- Given a set of points S in a space M , a query point $q \in M$, find the k closest points in S



Fixed Radius-NN

- Given a set of points S in a space M , a query point $q \in M$, find all the points in S , s.t., $\|s - q\| < r$





Why NN problem is important?



It is almost everywhere

- What we have covered:
 - Surface normal estimation
 - Noise filtering
 - Sampling
- What we will cover:
 - Clustering
 - Deep learning
 - Feature detection / description
 -



Why don't we simply call an open-source library (flann, PCL, etc.)?

- They are not efficient enough.
 - They are general lib, not optimized for 2D/3D.
 - Most open-source octree implementation is in-efficient, while octree is most effective for 3D.
- Few GPU based NN library is available



Why NN is difficult for point clouds



For Images, a neighbor is simply $x + \Delta x, y + \Delta y$



For point clouds

- Irregular – no grid based representation
- Curse of dimensionality
 - Non-trivial to build grids
 - Non-trivial to sort or build spatial partitions
- Huge data throughput in real-time applications
 - Velodyne HDL-64E – 2.2 million points per second / 110,000 points at 20Hz
 - Brute-force self-NN search is $110,000 \times 110,000 \times 0.5 = 6 \times 10^9$ comparisons @ 20Hz

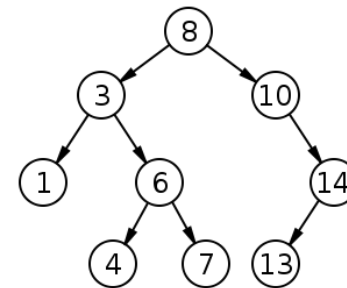


Lecture Outline



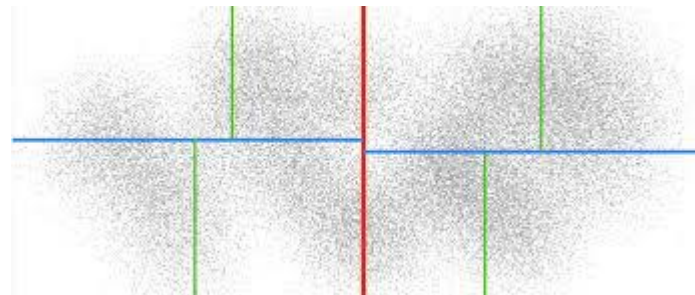
Binary Search Tree (BST)

- Basic knowledge about trees
- 1D NN problem
- With Python codes



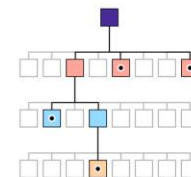
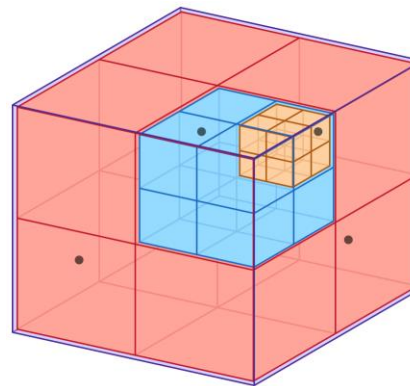
Kd-tree

- Works for data of any dimension
- Illustrated in 2D
- With Python codes



Octree

- Specifically designed for 3D data
- Illustrated in 2D/3D
- With Python codes





Core Ideas Shared by BST, kd-tree, octree



NN by space partition

- Split the space into different areas,
- Search some areas only, instead of all the data points



Stopping criteria

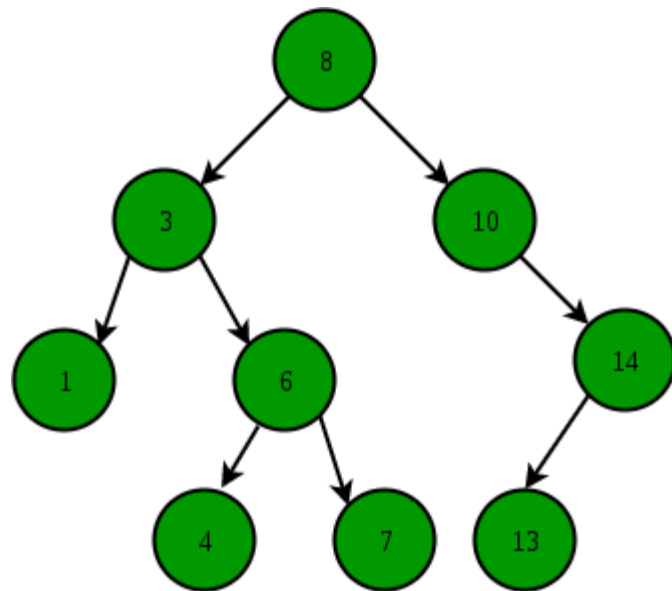
- How to skip some partitions?
 - Intersection of the “worst distance” with the partition boundaries
- How to stop the k-NN / Radius-NN search?
 - Search until the root
 - A partition completely contains the “worst distance”



Binary Search Tree (BST)

BST is a node-based tree data structure:

1. A node's left subtree contains nodes with keys lesser than its key
2. A node's right subtree contains nodes with keys larger than its key
3. The left / right subtree is BST





Binary Search Tree (BST)

From Wikipedia

Binary search tree

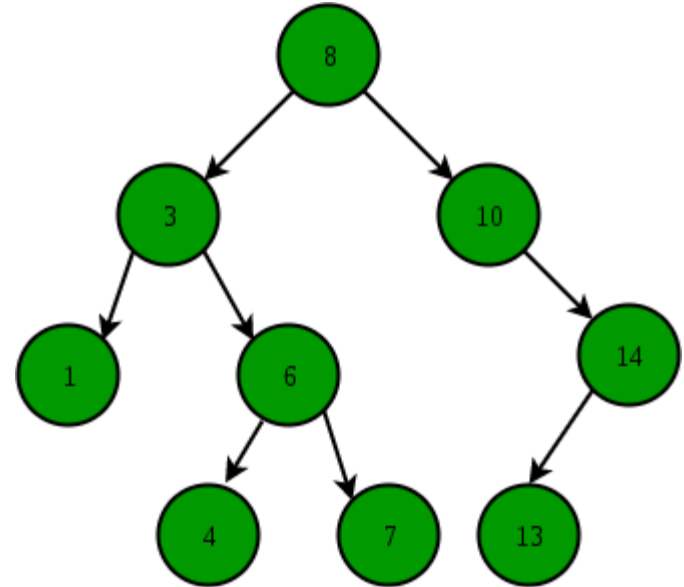
Type tree

Invented 1960

Invented by P.F. Windley, [A.D. Booth](#), [A.J.T. Colin](#), and [T.N. Hibbard](#)

Time complexity in big O notation

Algorithm	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(n)$
Insert	$O(\log n)$	$O(n)$
Delete	$O(\log n)$	$O(n)$





BST – Node definition



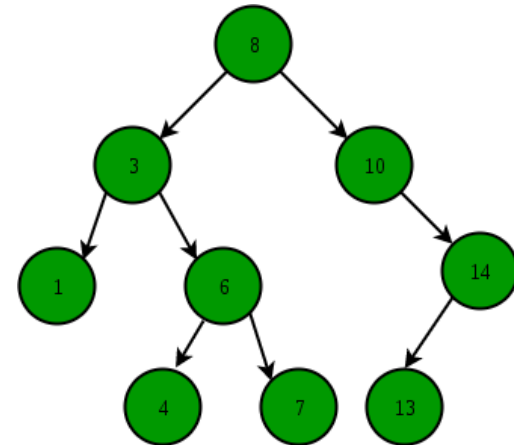
A node contains

1. Key
2. Left child
3. Right child
4.



The left/right child is a Node as well

```
class Node:
    def __init__(self, key, value=-1):
        self.left = None
        self.right = None
        self.key = key
        self.value = value
```





BST – Construction / Insertion



Given N 1D-points (scalar) denoted by an array

$$\{x_1, x_2, \dots, x_n\}, x_i \in \mathbb{R}$$



Construct a BST that stores the points and its index in the array, e.g.

[100, 20, 500, 10, 30, 40]

Data generation

```
db_size = 10  
data = np.random.permutation(db_size).tolist()
```

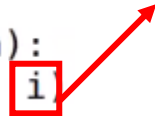
Recursively insert each an element

```
def insert(root, key, value=-1):  
    if root is None:  
        root = Node(key, value)  
    else:  
        if key < root.key:  
            root.left = insert(root.left, key, value)  
        elif key > root.key:  
            root.right = insert(root.right, key, value)  
        else: # don't insert if key already exist in the tree  
            pass  
    return root
```

Insert each element

```
root = None  
for i, point in enumerate(data):  
    root = insert(root, point, i)
```

“value” in the Node is the index of a point in the array
Useful in later NN search





BST – Insertion Complexity



The worst case is $O(h)$, where h is the height of the BST



In the worst case, h is the number of points in BST.

- Unbalanced tree – a chain in an extreme case
- E.g., inserting [9, 8, 7, 6, 5, 4, 3] into an empty BST

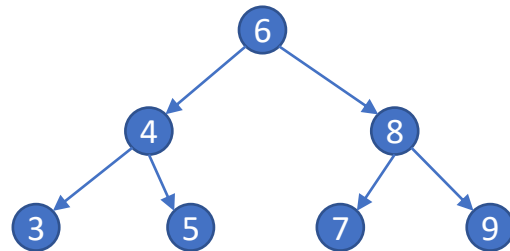
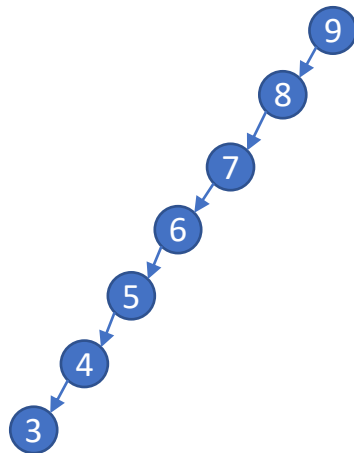


Tree balancing is another topic

- Sort the array and insert as a balanced tree (select median as root)
- AVL tree
- Red-Black tree
- etc.



Best case $h = \log_2 n$





BST – Search



Given a BST, and a query (key), determine *which node* equals to that query (key), if not, return *NULL*



Can be done *recursively* or *iteratively*



BST – Search Recursively

Search till the leaf but not found.

Find a match

For sure just need to look at the left

```
def search_recursive(root, key):  
    if root is None or root.key == key:  
        return root  
    if key < root.key:  
        return search_recursive(root.left, key)  
    elif key > root.key:  
        return search_recursive(root.right, key)
```

For sure just need to look at the right



BST – Search Iteratively



Use “current_node” to simulate a *Stack*, so that recursion is avoided



In any case, you can write your own *Stack* to avoid recursion, but that may be complicated sometimes.

```
def search_iterative(root, key):
    current_node = root
    while current_node is not None:
        if current_node.key == key:
            return current_node
        elif key < current_node.key:
            current_node = current_node.left
        elif key > current_node.key:
            current_node = current_node.right
    return current_node
```



Search recursively or iteratively complexity same as insertion
worst $O(h)$



BST – Search

Recursion

Pros:

- Easy to understand / implement
- Codes are short

Cons:

- Hard to trace step-by-step
- $O(n)$ storage, n is number of recursion (May be optimized by compiler)

Iteration

Pros:

- Avoid [stack-overflow](#), e.g., in embedded system / GPU
- Easier in step-by-step tracing
- $O(1)$ storage

Cons:

- The logic is complicated



BST – Depth First Traversal



Inorder – Left, Root, Right

- E.g., sorting
- 1, 3, 4, 6, 7, 8, 10, 13, 14



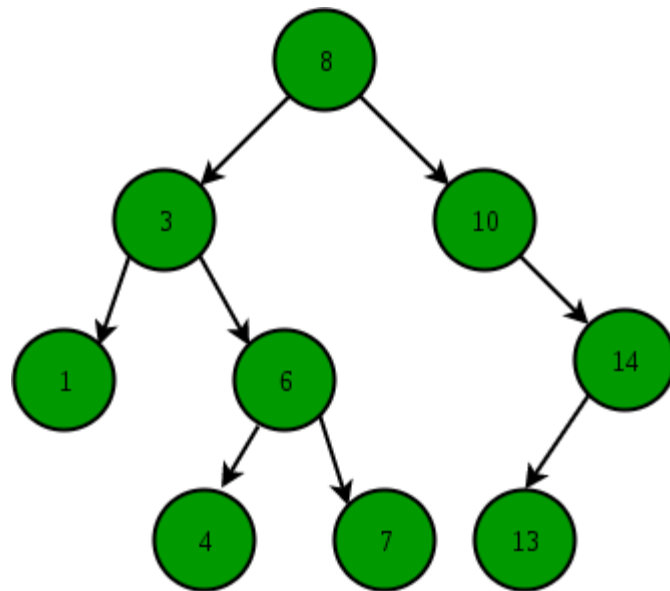
Preorder – Root, Left, Right

- E.g., copy a tree
- 8, 3, 1, 6, 4, 7, 10, 14, 13



Postorder – Left, Right, Root

- E.g., delete a node
- 1, 4, 7, 6, 3, 13, 14, 10, 8



```
def inorder(root):  
    # Inorder (Left, Root, Right)  
    if root is not None:  
        inorder(root.left)  
        print(root)  
        inorder(root.right)
```

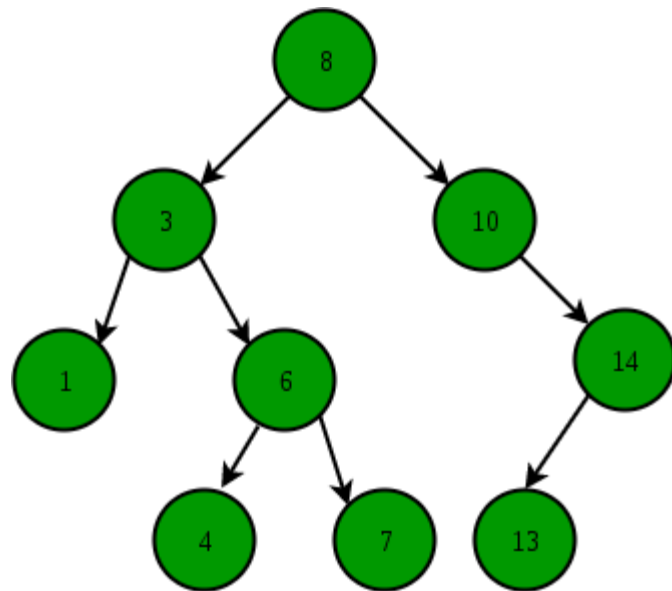
1, 3, 4, 6, 7, 8, 10, 13, 14

```
def preorder(root):  
    # Preorder (Root, Left, Right)  
    if root is not None:  
        print(root)  
        preorder(root.left)  
        preorder(root.right)
```

8, 3, 1, 6, 4, 7, 10, 14, 13

```
def postorder(root):  
    # Postorder (Left, Right, Root)  
    if root is not None:  
        postorder(root.left)  
        postorder(root.right)  
        print(root)
```

1, 4, 7, 6, 3, 13, 14, 10, 8

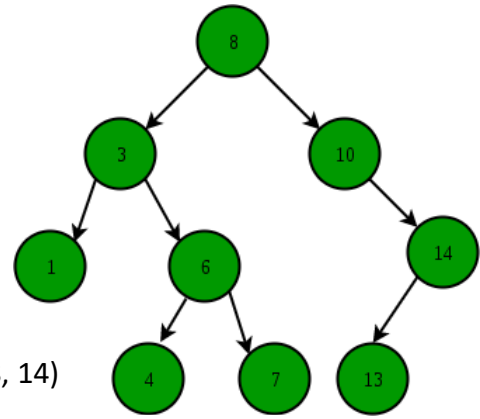




BST – 1NN Search

- Query point – 11

1. 8
 - a) **worst distance = 3** (11-8)
 - b) any point in 8's left tree is at least 3 away from query
 - c) do I need to go further? Yes! Right subtree is (8, +inf] but **worst distance=3** -> need to search (8, 14)
- 2.
- 3.
- 4.
- 5.
- 6.





BST – kNN Search



Almost same as 1NN search



Difference is how to compute **worst distance**



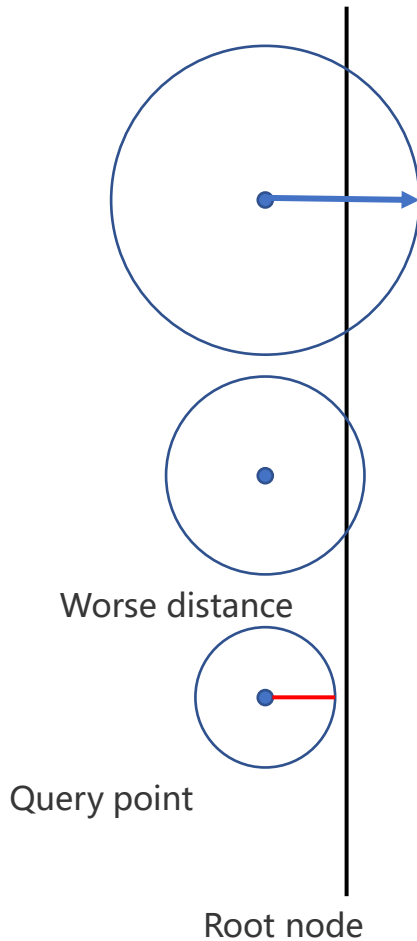
Worst distance is the largest distance that you should search around the query point



Areas outside this “worst circle” can be skipped



In kNN search, the worst distance is dynamic





BST – Worst Distance for kNN



Build a container to store the kNN results



k results are sorted



worst_dist is the last one



Add a result if

$$dist < worse_dist$$

- Example:
- Existing container content
- [1, 2, 3, 4, inf, inf]
- **add_point(3.5)**
- *Step 1.* Make space for 3.5
- [1, 2, 3, 4, 4, inf]
- *Step 2.* Put 3.5 in the correct position
- [1, 2, 3, 3.5, 4, inf]
- *Step 3.* Update worst_dist

```
class KNNResultSet:
```

```
def __init__(self, capacity):
    self.capacity = capacity
    self.count = 0
    self.worst_dist = 1e10
    self.dist_index_list = []
    for i in range(capacity):
        self.dist_index_list.append(DistIndex(self.worst_dist, 0))
    self.comparison_counter = 0
```

Initialized to large value

Container to keep all the k neighbors

```
def size(self):
    return self.count

def full(self):
    return self.count == self.capacity
```

```
def worstDist(self):
    return self.worst_dist
```

If a point is added, put it in a ordered position

```
def add_point(self, dist, index):
    self.comparison_counter += 1
    if dist > self.worst_dist:
        return

    if self.count < self.capacity:
        self.count += 1

    i = self.count - 1
    while i > 0:
        if self.dist_index_list[i-1].distance > dist:
            self.dist_index_list[i] = copy.deepcopy(self.dist_index_list[i-1])
            i -= 1
        else:
            break

    self.dist_index_list[i].distance = dist
    self.dist_index_list[i].index = index
    self.worst_dist = self.dist_index_list[self.capacity-1].distance
```

```
class DistIndex:
```

```
def __init__(self, distance, index):
    self.distance = distance
    self.index = index

def __lt__(self, other):
    return self.distance < other.distance
```

```

def knn_search(root: Node, result_set: KNNResultSet, key):
    if root is None:
        return False

    # compare the root itself
    result_set.add_point(math.fabs(root.key - key), root.value)
    if result_set.worstDist() == 0:
        return True

    if root.key >= key:
        # iterate left branch first
        if knn_search(root.left, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return knn_search(root.right, result_set, key)
        return False
    else:
        # iterate right branch first
        if knn_search(root.right, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return knn_search(root.left, result_set, key)
        return False

```

A special case – if the worst distance is 0, no need to search anymore

If key != query, need to go through one subtree

May not need to search for the other subtree, depends on worst distance

Similar to the “if” block above



Radius NN Search

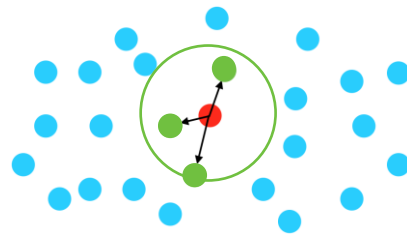
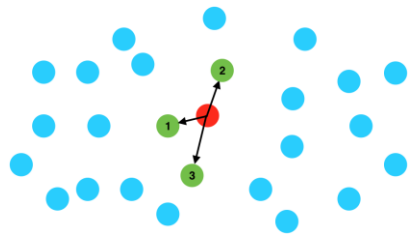


Same as kNN, in the sense that,

- Worst distance circle intersects the boundary -> search
- If not -> skip



In implementation, we don't need to change the BST kNN search logics, except that, the worst distance is **fixed**, instead of dynamic





BST – Radius Result Set Manger



Simpler than kNN result set manager.



Worst distance is fixed.



No need to maintain a sorted result set.

```
def add_point(self, dist, index):  
    self.comparison_counter += 1  
    if dist > self.radius:  
        return  
  
    self.count += 1  
    self.dist_index_list.append(DistIndex(dist, index))
```



BST - kNN v.s. Radius NN

```
def knn_search(root: Node, result_set: KNNResultSet, key):
    if root is None:
        return False

    # compare the root itself
    result_set.add_point(math.fabs(root.key - key), root.value)
    if result_set.worstDist() == 0:
        return True

    if root.key >= key:
        # iterate left branch first
        if knn_search(root.left, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return knn_search(root.right, result_set, key)
        return False
    else:
        # iterate right branch first
        if knn_search(root.right, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return knn_search(root.left, result_set, key)
        return False
```

This part is gone in
radius search, because
worst_dist = r

```
def radius_search(root: Node, result_set: RadiusNNResultSet, key):
    if root is None:
        return False

    # compare the root itself
    result_set.add_point(math.fabs(root.key - key), root.value)

    if root.key >= key:
        # iterate left branch first
        if radius_search(root.left, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return radius_search(root.right, result_set, key)
        return False
    else:
        # iterate right branch first
        if radius_search(root.right, result_set, key):
            return True
        elif math.fabs(root.key-key) < result_set.worstDist():
            return radius_search(root.left, result_set, key)
        return False
```



A complete script

```
db_size = 100
k = 5
radius = 2.0

data = np.random.permutation(db_size).tolist()

root = None
for i, point in enumerate(data):
    root = insert(root, point, i)

query_key = 6
result_set = KNNResultSet(capacity=k)
knn_search(root, result_set, query_key)
print('kNN Search:')
print('index - distance')
print(result_set)

result_set = RadiusNNResultSet(radius=radius)
radius_search(root, result_set, query_key)
print('Radius NN Search:')
print('index - distance')
print(result_set)
```

- Search in 100 points, takes 7 comparison only
- Complexity is around $O(\log_2(n))$, n is number of database points, if tree is balanced
- Worst $O(N)$

kNN Search:

index - distance

73 - 0.00

5 - 1.00

12 - 1.00

1 - 2.00

98 - 2.00

In total 7 comparison operations.

Radius NN Search:

index - distance

73 - 0.00

5 - 1.00

12 - 1.00

1 - 2.00

98 - 2.00

In total 5 neighbors within 2.000000.

There are 7 comparison operations.

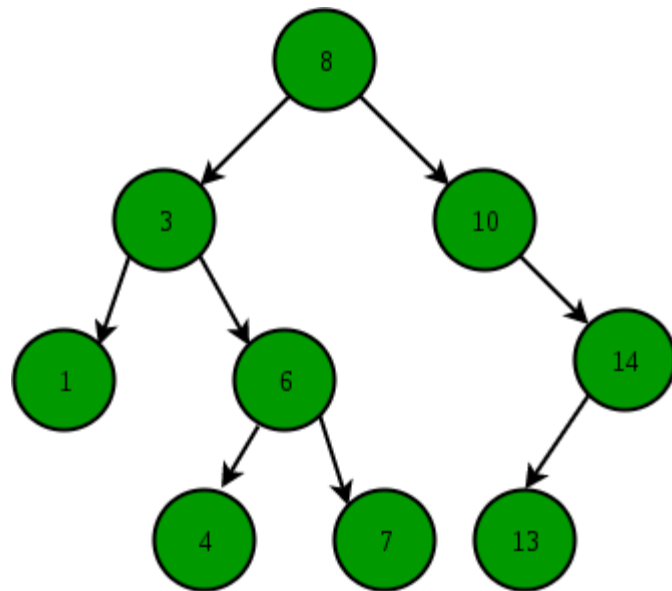


Binary Search Tree (BST)

BST based 1D kNN/RadiusNN search

- Naïve BST is for 1D data only

Tree based kNN/RadiusNN can be viewed as a Branch-n-Bound algorithm.





Kd-tree (k-dimensional tree)



It is an extension of BST into high dimension

- BST is 1-dimensional, how to extend?
- BST in each dimension!

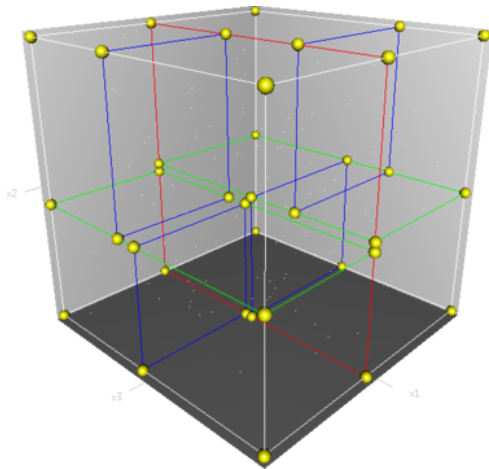


Invented by Jon Louis Bentley in 1975



The kd-tree is a binary tree where every leaf node is a **k-d**imensional point

k-d tree		
Type	Multidimensional BST	
Invented	1975	
Invented by	Jon Louis Bentley	
Time complexity in big O notation		
Algorithm	Average	Worst case
Space	$O(n)$	$O(n)$
Search	$O(\log n)$	$O(n)$
Insert	$O(\log n)$	$O(n)$
Delete	$O(\log n)$	$O(n)$



A 3-dimensional kd tree:

1. Red
2. Green
3. Blue



Kd-tree Construction



If there is only one point, or number of points $<$ leaf_size, build a leaf



Otherwise, divide the points in half by a hyperplane perpendicular to the selected splitting axis

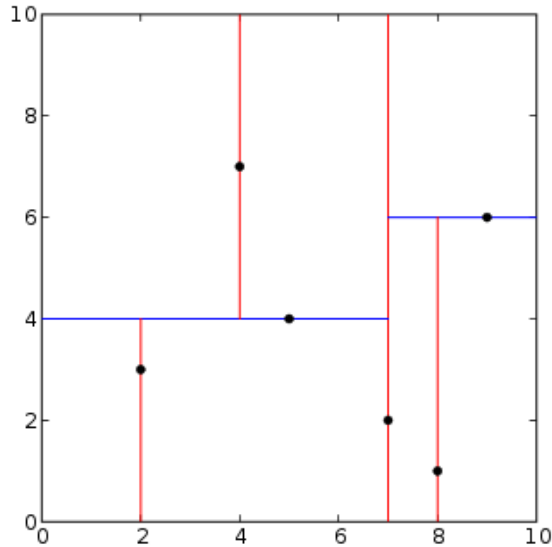


Recursively repeat the first two steps.

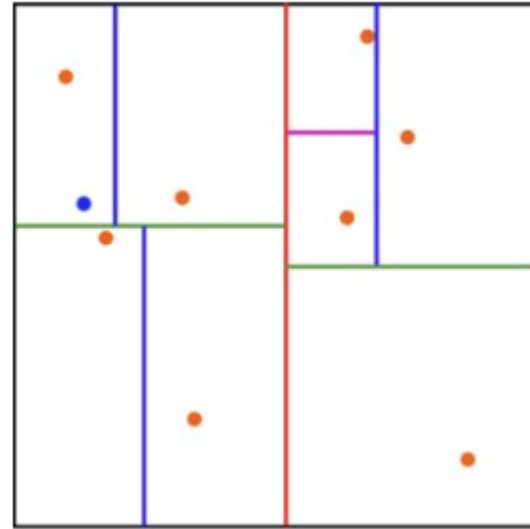


Kd-tree Construction – Two Conventions

Splitting position is one of the points



Splitting position is **NOT** one of the points

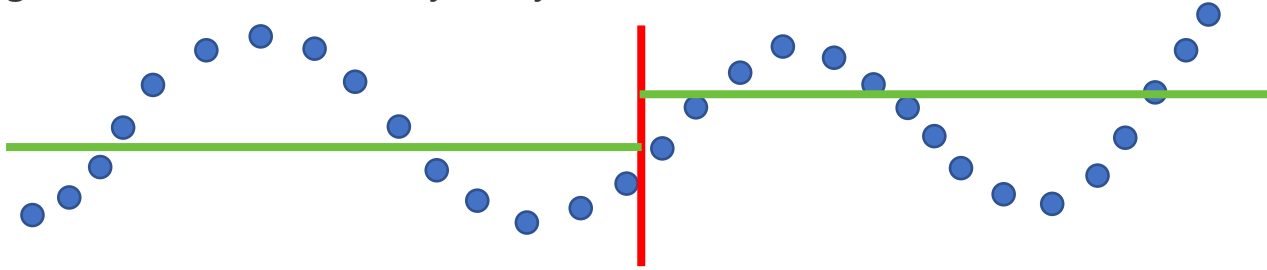




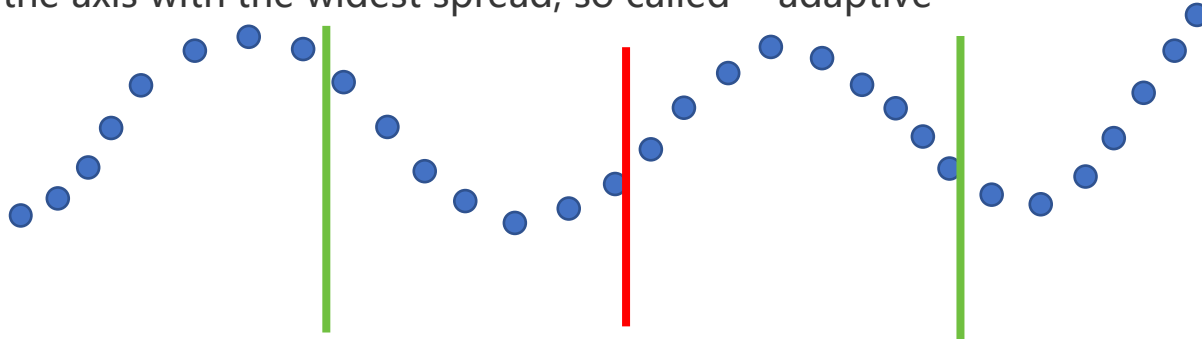
Kd-tree Construction

Division / Splitting Strategy

- Dividing axis is round-robin: x-y-z-x-y-z-x-.....

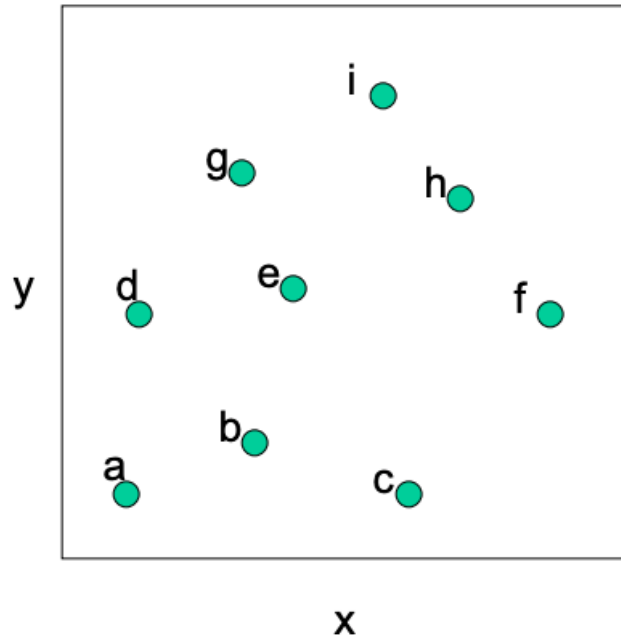


- Select the axis with the widest spread, so called "adaptive"



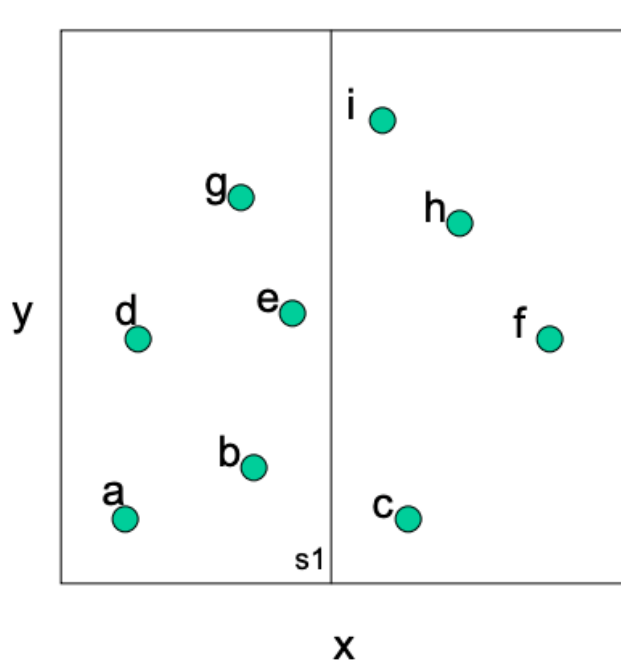


Kd-tree Construction





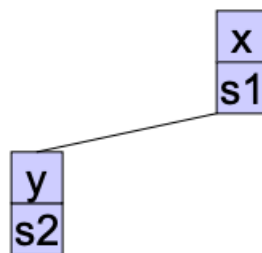
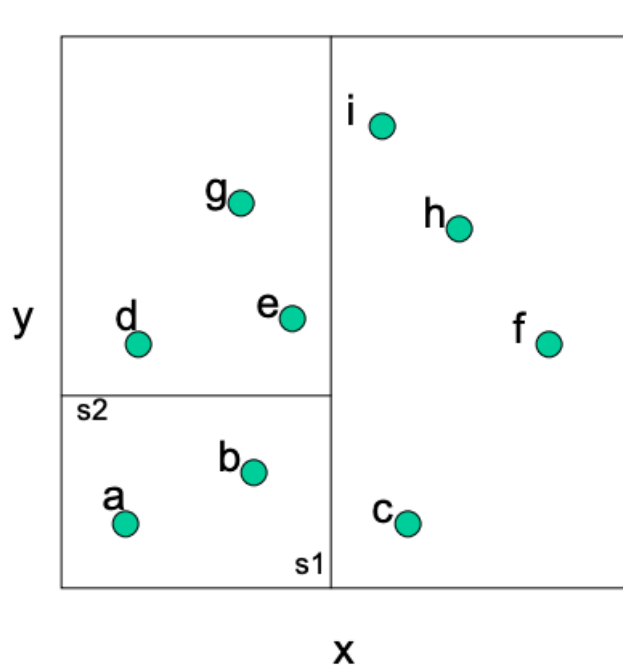
Kd-tree Construction



x
s1

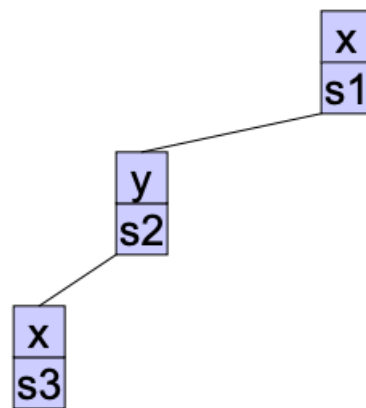
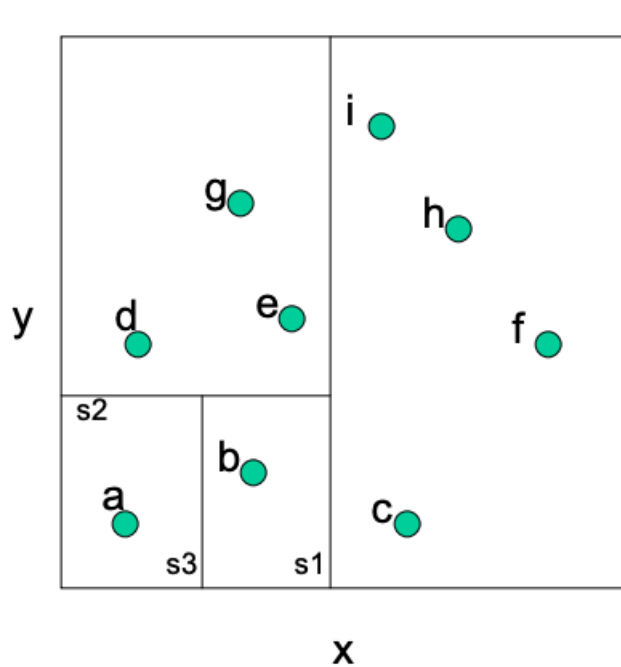


Kd-tree Construction



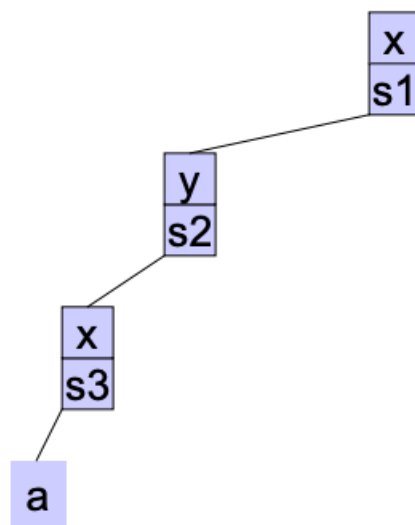
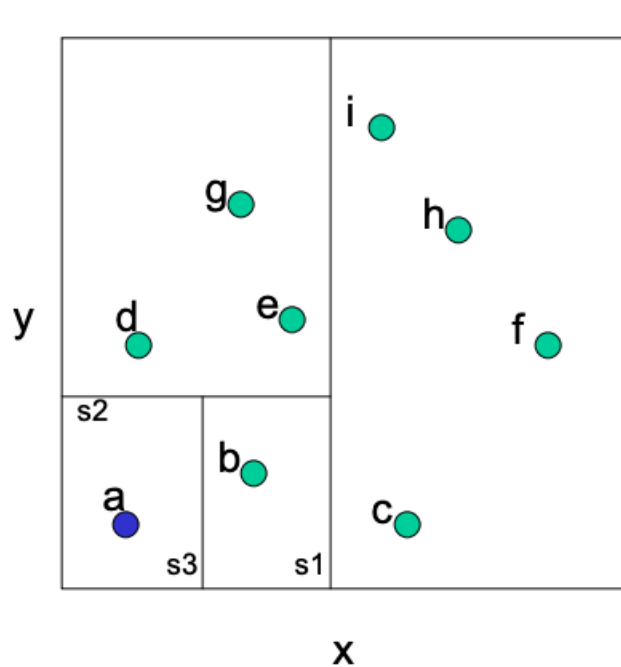


Kd-tree Construction



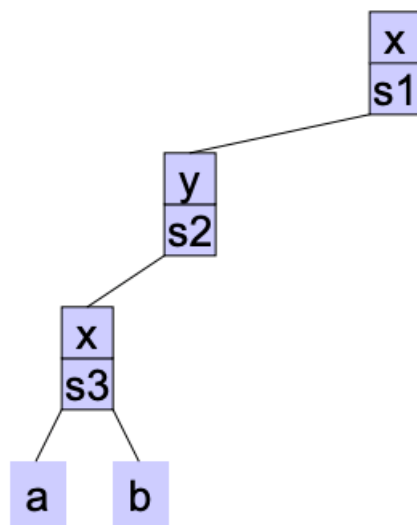
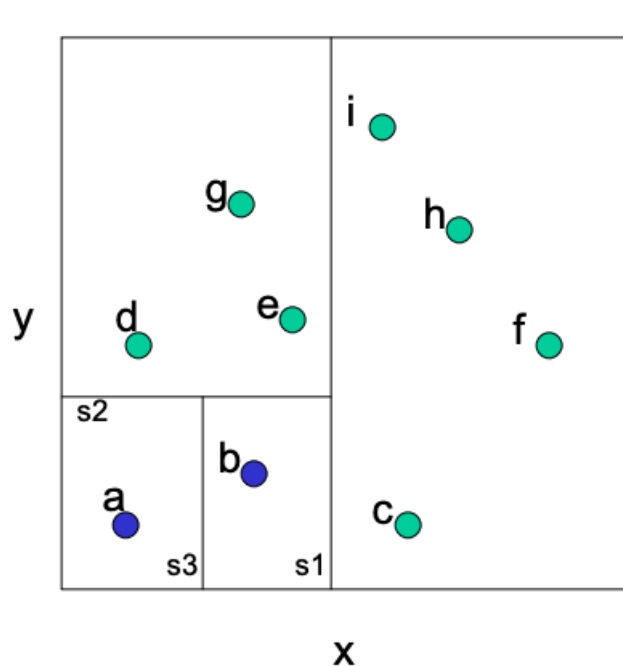


Kd-tree Construction



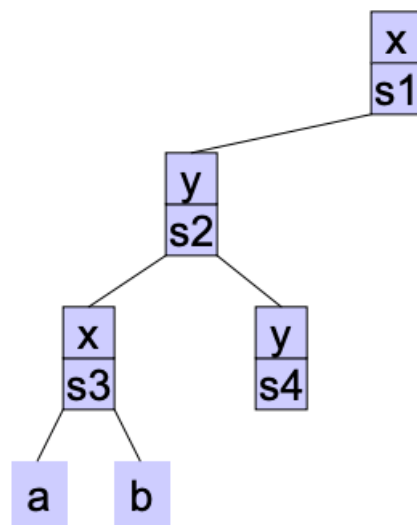
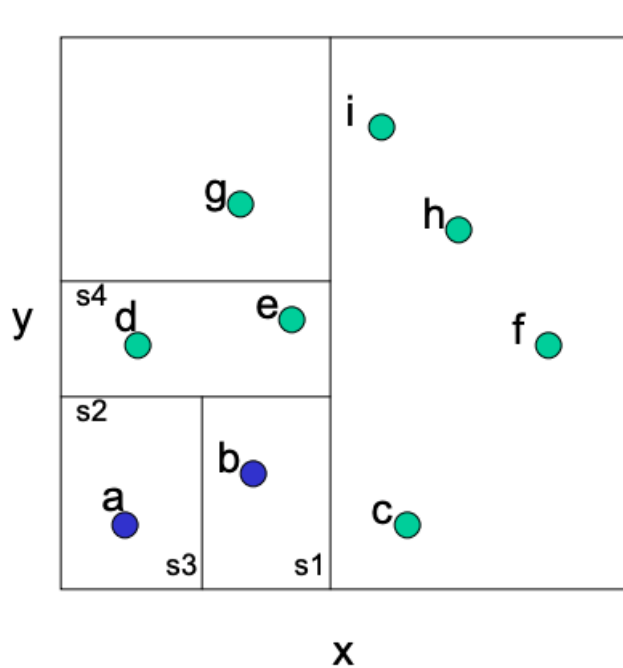


Kd-tree Construction



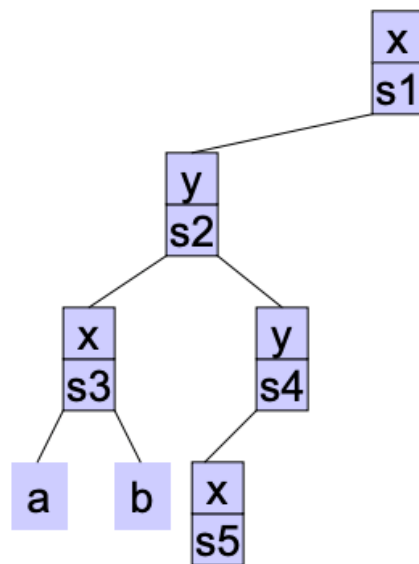
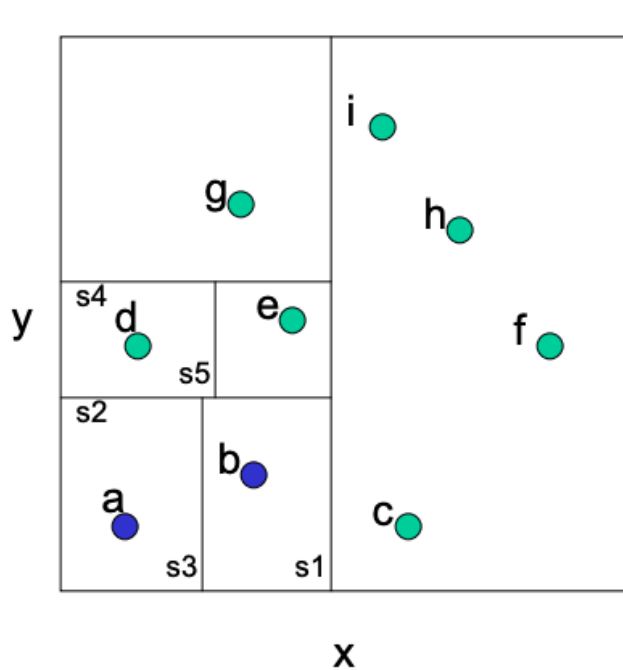


Kd-tree Construction



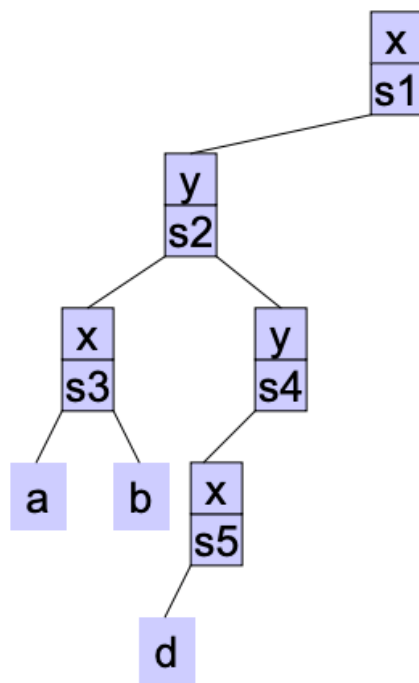
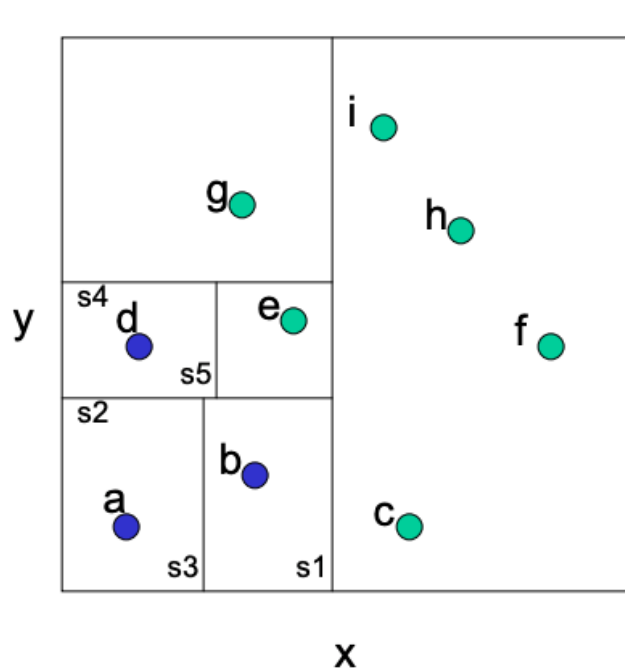


Kd-tree Construction



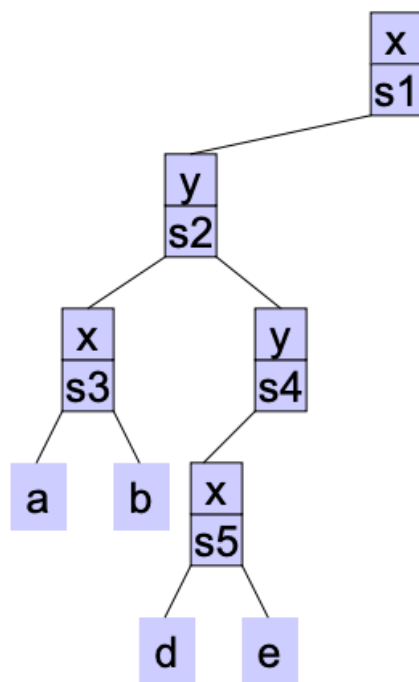
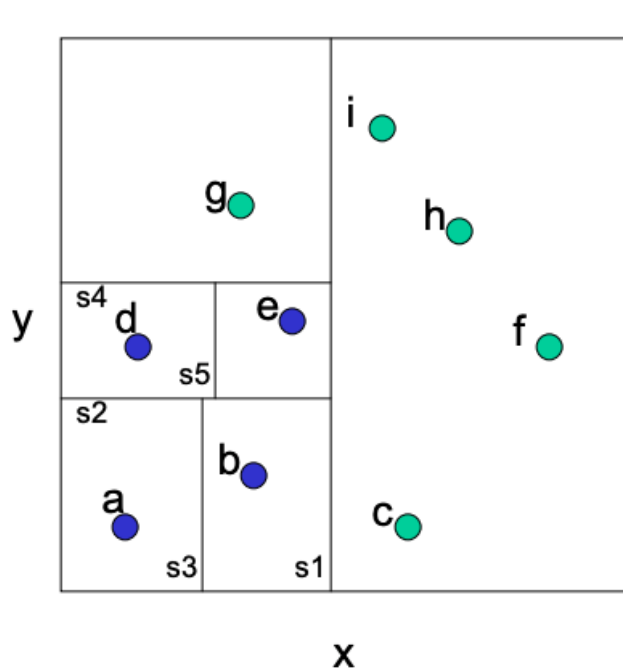


Kd-tree Construction



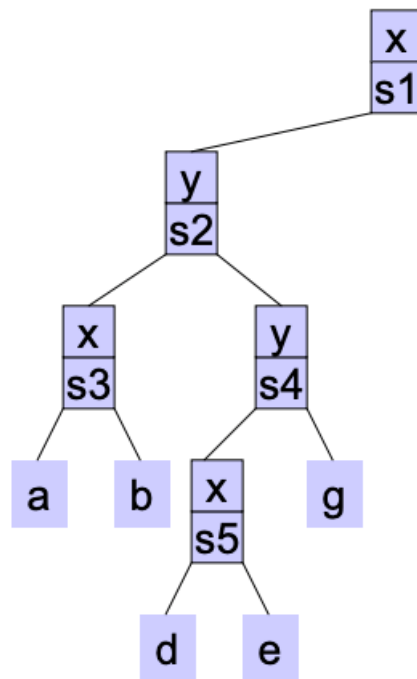
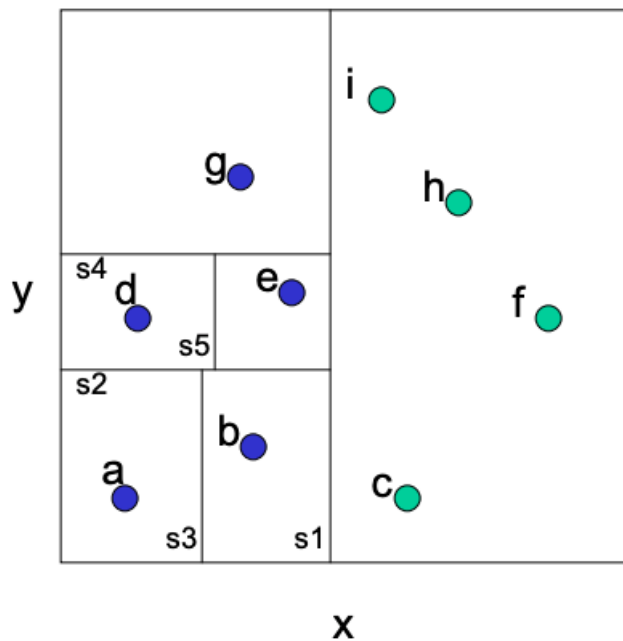


Kd-tree Construction



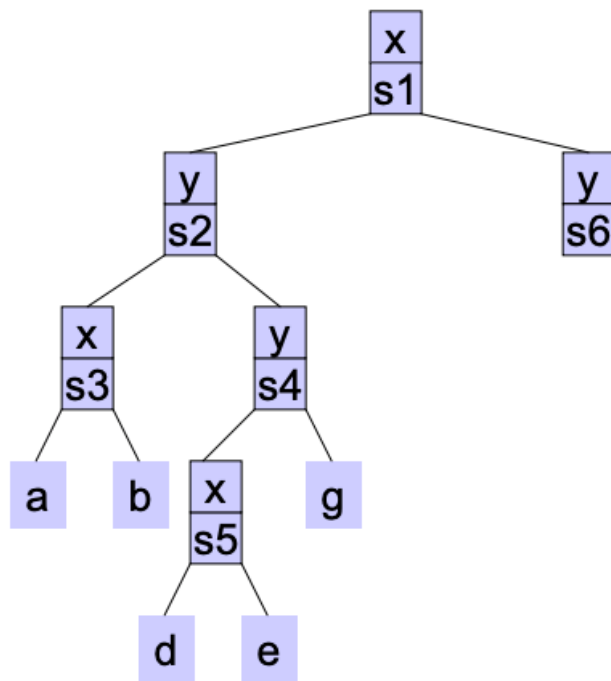
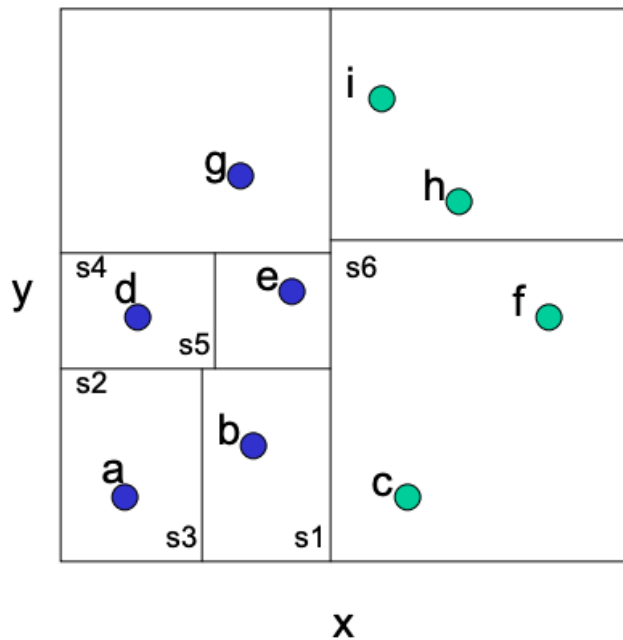


Kd-tree Construction



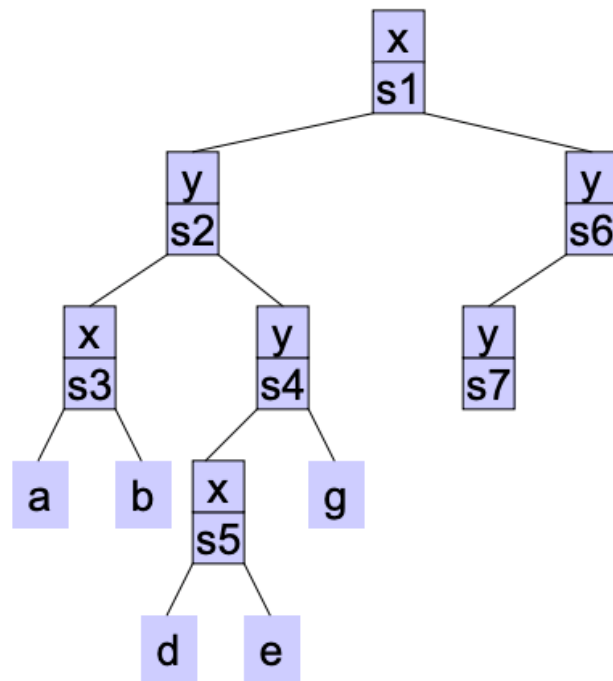
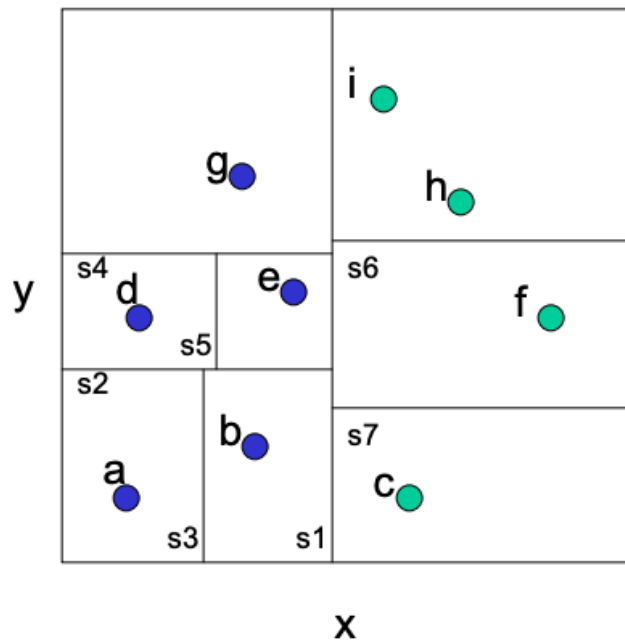


Kd-tree Construction



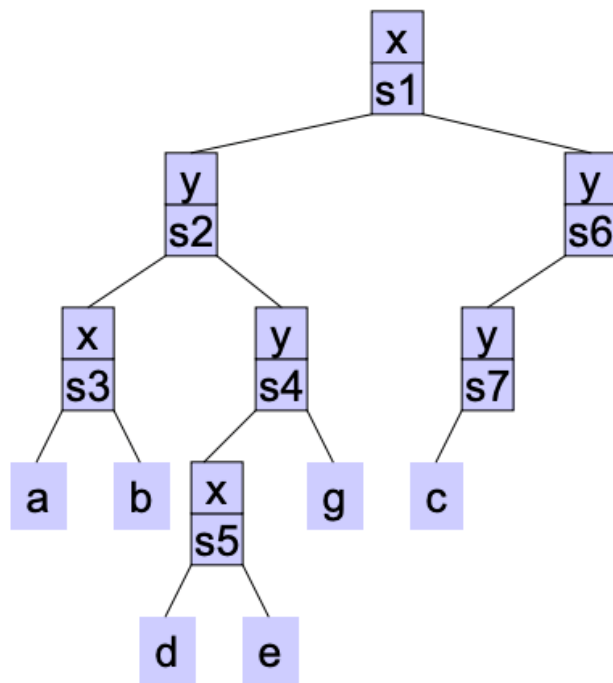
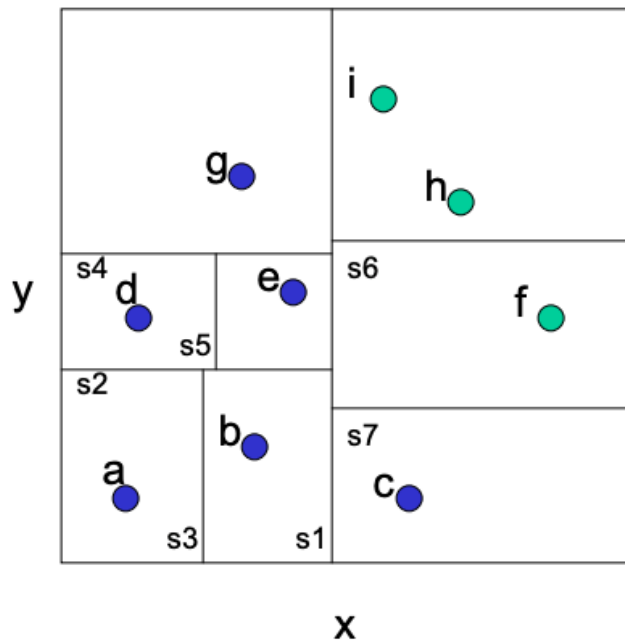


Kd-tree Construction



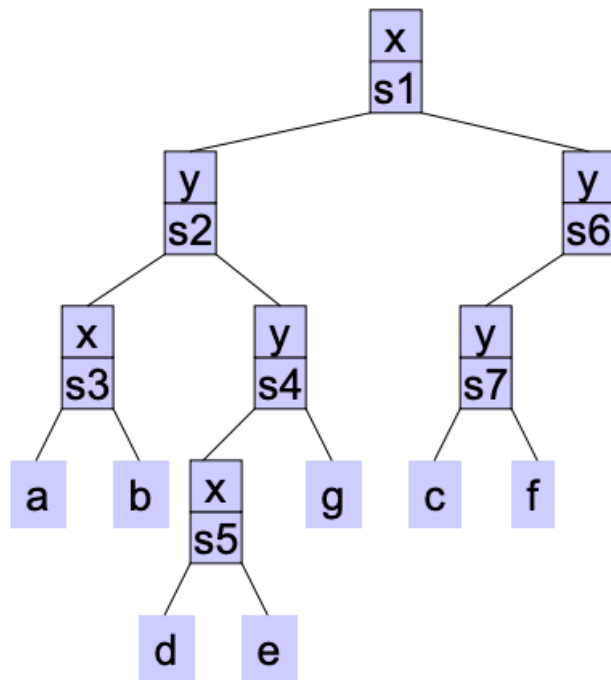
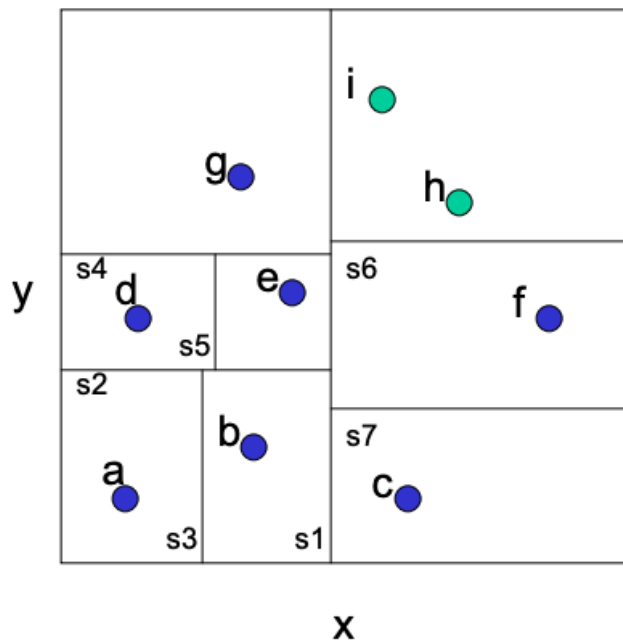


Kd-tree Construction



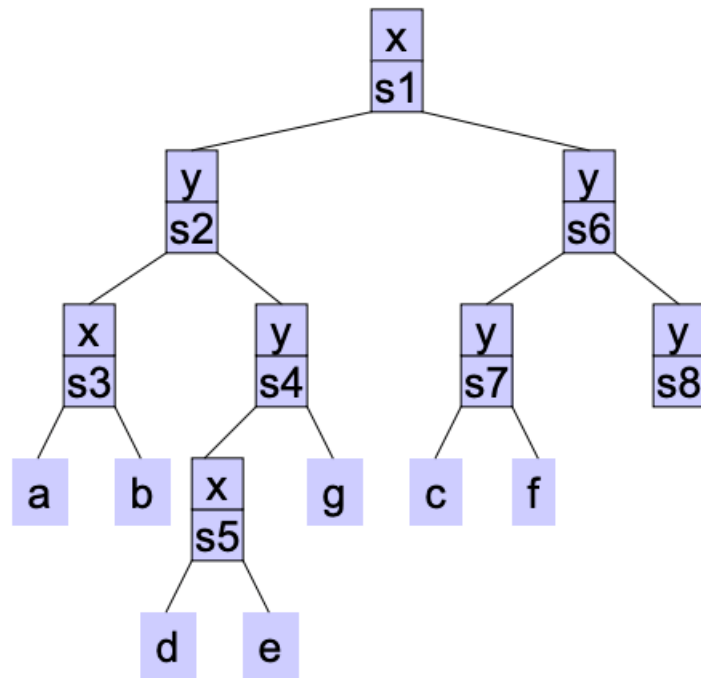
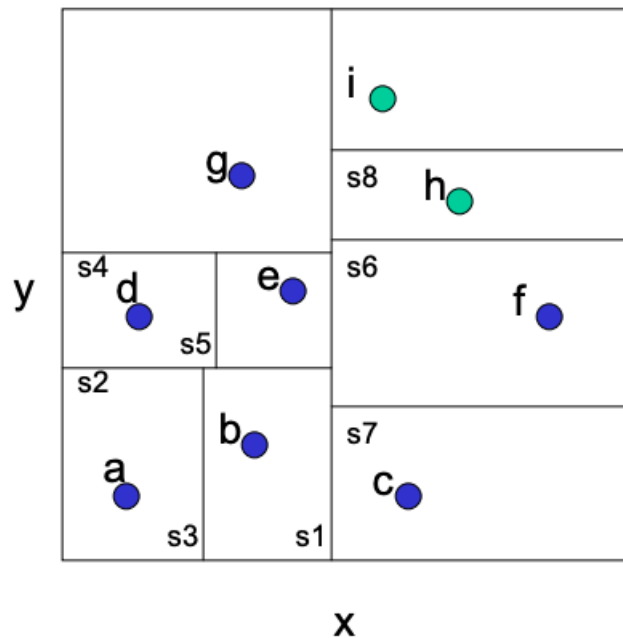


Kd-tree Construction



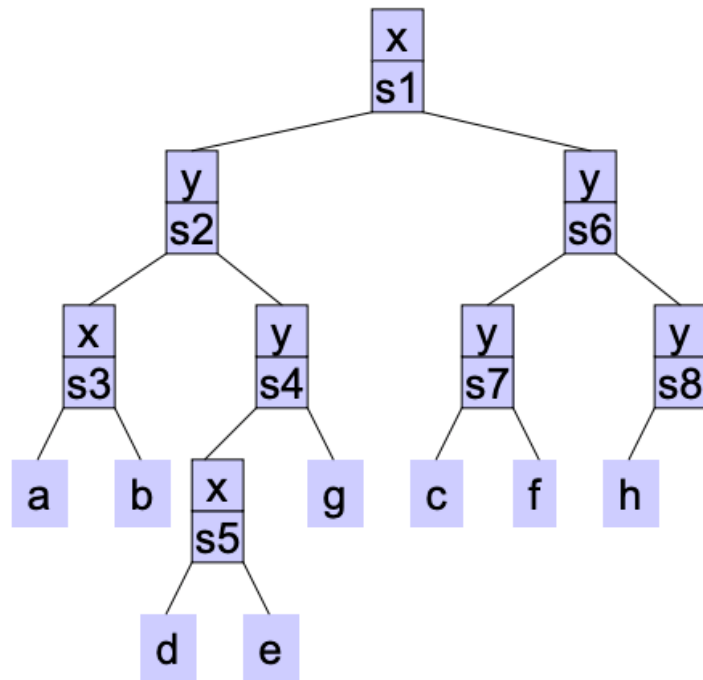
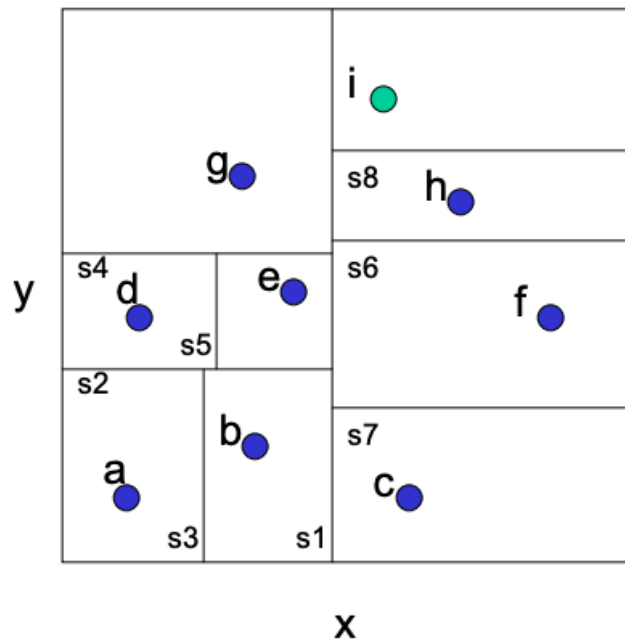


Kd-tree Construction



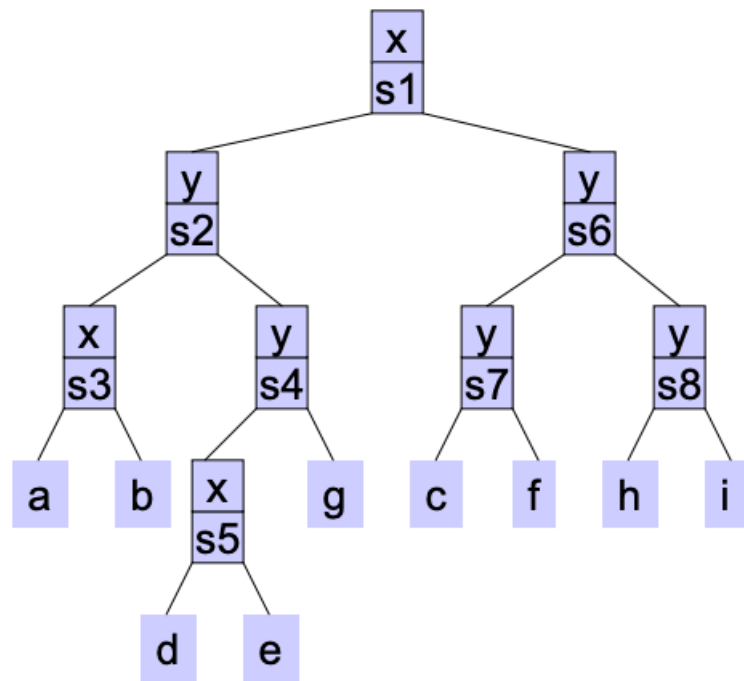
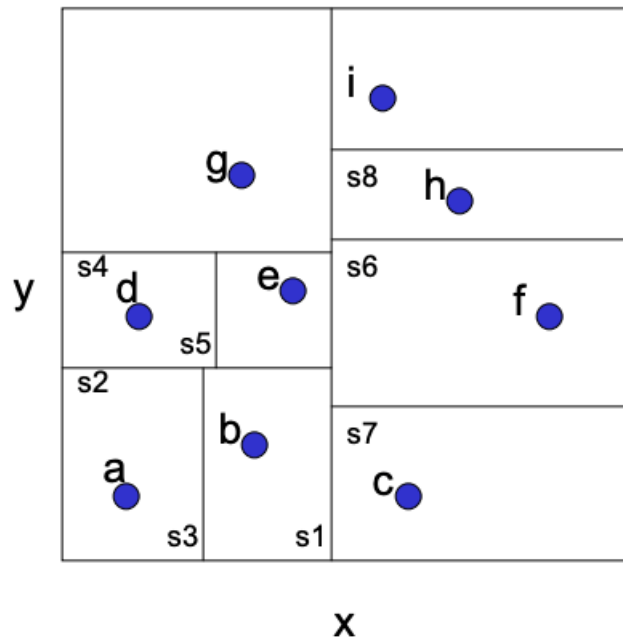


Kd-tree Construction





Kd-tree Construction





Kd-tree Construction



Talk is cheap, show me the code.

Linus Torvalds



Kd-tree Node Representation

```
class Node:
    def __init__(self, axis, value, left, right, point_indices):
        self.axis = axis
        self.value = value
        self.left = left
        self.right = right
        self.point_indices = point_indices

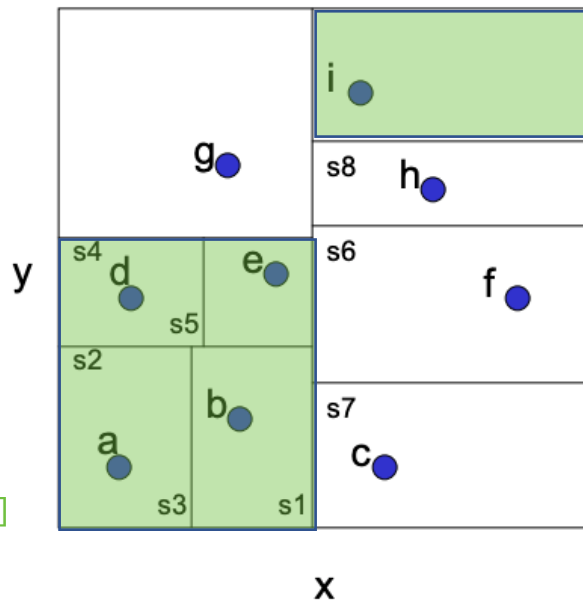
    def is_leaf(self):
        if self.value is None:
            return True
        else:
            return False
```

Splitting position

Stores points that belongs to this partition

A node with
axis = x
value = ***
points = [a, b, d, e]

A leaf node with
axis = y
value = None
points = i



```
def kdtree_recursive_build(root, db, point_indices, axis, leaf_size):
```

A leaf node can contain
more than 1 point

```
    :param root:
    :param db: NxM
    :param db_sorted_idx_inv: NxM
    :param point_idx: M
    :param axis: scalar
    :param leaf_size: scalar
    :return:
```

```
    """
    if root is None:
        root = Node(axis, None, None, None, point_indices)
```

```
    # determine whether to split into left and right
```

```
    if len(point_indices) > leaf_size:
```

```
        # --- get the split position ---
```

```
        point_indices_sorted, _ = sort_key_by_val(point_indices, db[point_indices, axis]) # M
```

```
        middle_left_idx = math.ceil(point_indices_sorted.shape[0] / 2) - 1
```

```
        middle_left_point_idx = point_indices_sorted[middle_left_idx]
```

```
        middle_left_point_value = db[middle_left_point_idx, axis]
```

```
        middle_right_idx = middle_left_idx + 1
```

```
        middle_right_point_idx = point_indices_sorted[middle_right_idx]
```

```
        middle_right_point_value = db[middle_right_point_idx, axis]
```

```
        root.value = (middle_left_point_value + middle_right_point_value) * 0.5
```

```
        # === get the split position ===
```

```
        root.left = kdtree_recursive_build(root.left,
```

```
            db,
```

```
            point_indices_sorted[0:middle_right_idx],
```

```
            axis_round_robin(axis, dim=db.shape[1]),
```

```
            leaf_size)
```

```
        root.right = kdtree_recursive_build(root.right,
```

```
            db,
```

```
            point_indices_sorted[middle_right_idx:],
```

```
            axis_round_robin(axis, dim=db.shape[1]),
```

```
            leaf_size)
```

```
    return root
```

Sort the points in this node, get the
median position

```
def axis_round_robin(axis, dim):
    if axis == dim-1:
        return 0
    else:
        return axis + 1
```



Kd-tree Construction Complexity



The example shown here is not optimal because of sorting at each level of the tree

- Time complexity of around $O(n \log n \log n)$
- Space complexity of $O(kn + n \log n) \rightarrow$ can be easily reduced to $O(kn + n)$
 - Only store points at leaf



Can we select median instead of sorting?

- If median finding is $O(n)$
- Kd-tree is $O(n \log n)$
- Median finding in $O(n)$ is complicated, but **possible!**



$O(kn \log n)$ method

- Building a Balanced k-d Tree in $O(kn \log n)$ Time
 - Russel A. Brown, Journal of Computer Graphics Techniques, 2015



Kd-tree Construction Complexity



Simple methods that work well in practice

- Sample a **subset** of point in each node for sorting, instead of sorting all points
 - $O(n' \log n)$ or $O(n' \log n' \log n)$
- Use **mean** instead of median
 - An easy way to achieve $O(n \log n)$



They don't guarantee balanced kd-tree

- Balanced tree - each leaf node is approximately the same distance from the root
- Similar to the "chain" example in BST.



Kd-tree – kNN Search



Start from root



Reach the leaf node than covers the query point

- Compare all points in the leaf node

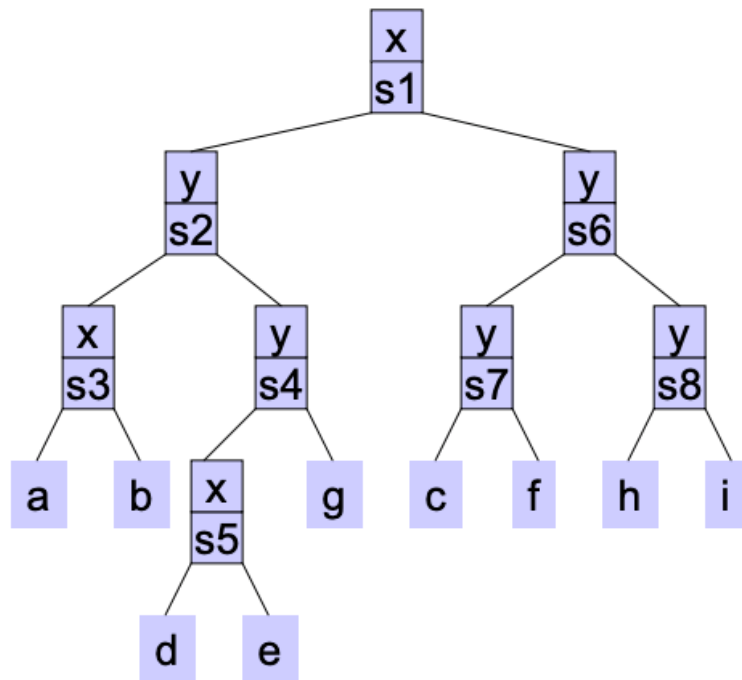
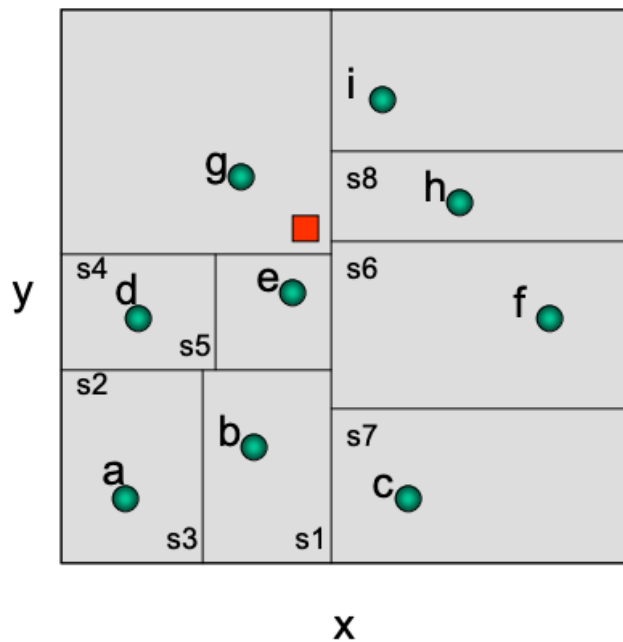


Go up and traverse the tree



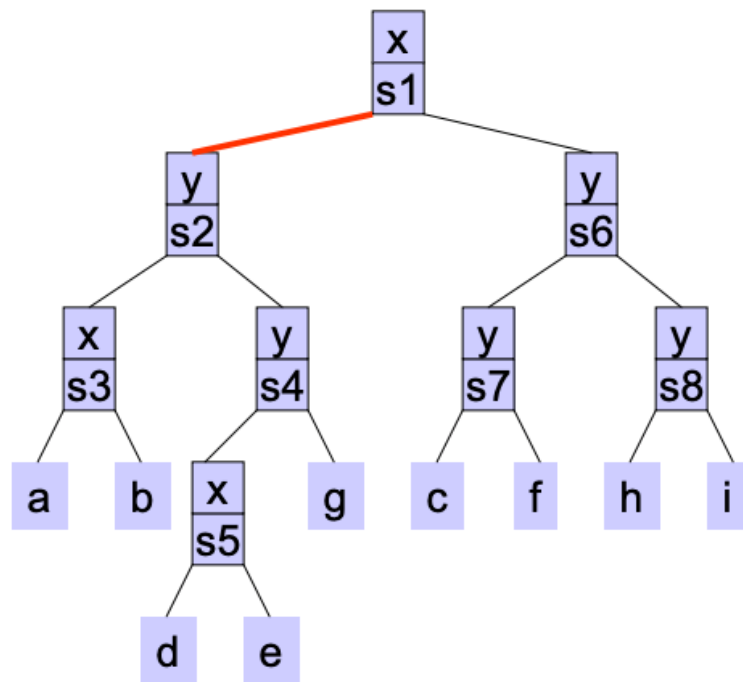
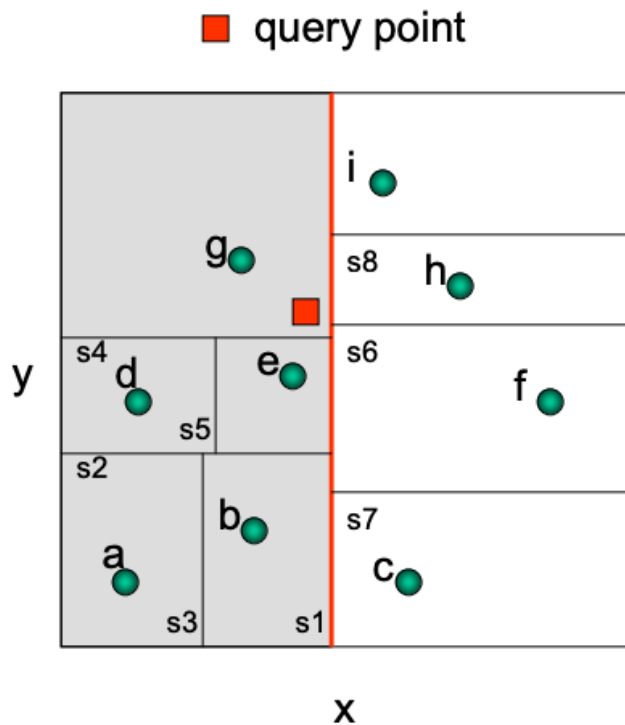
Kd-tree – kNN Search

■ query point



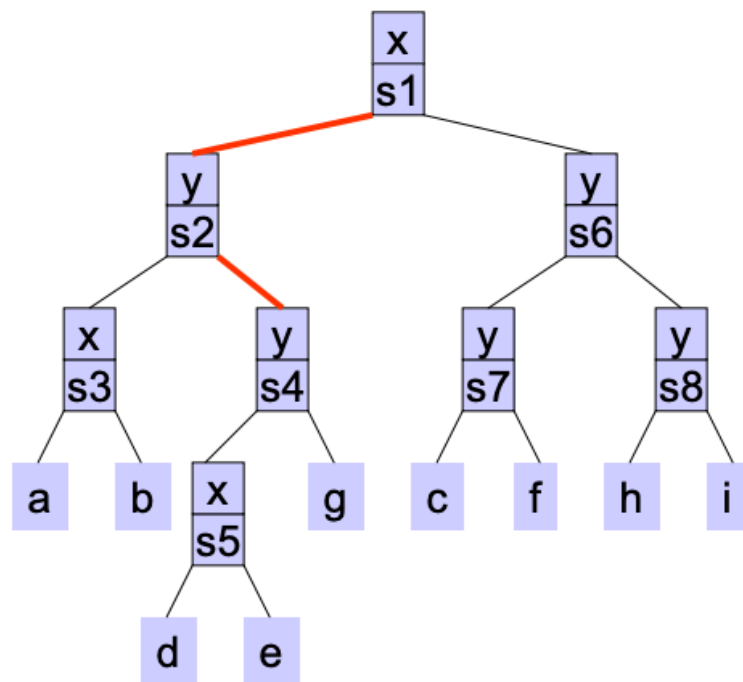
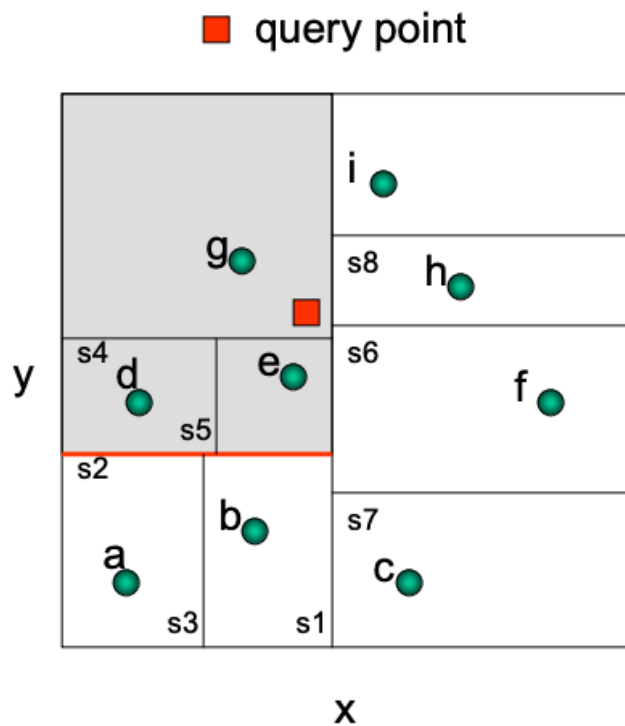


Kd-tree – kNN Search



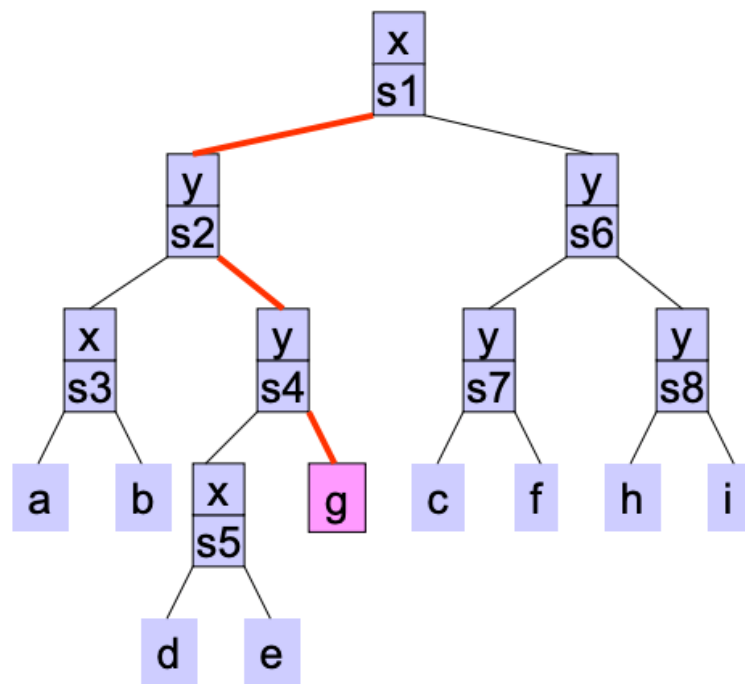
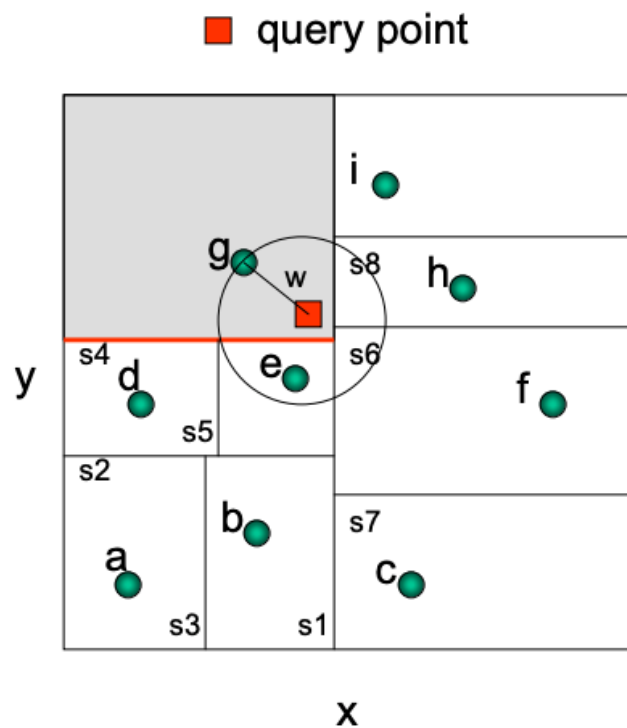


Kd-tree – kNN Search



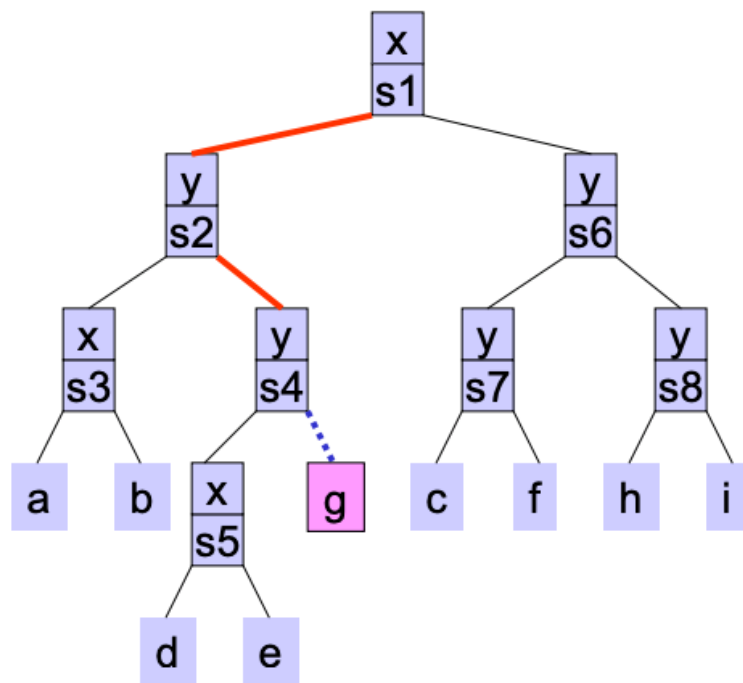
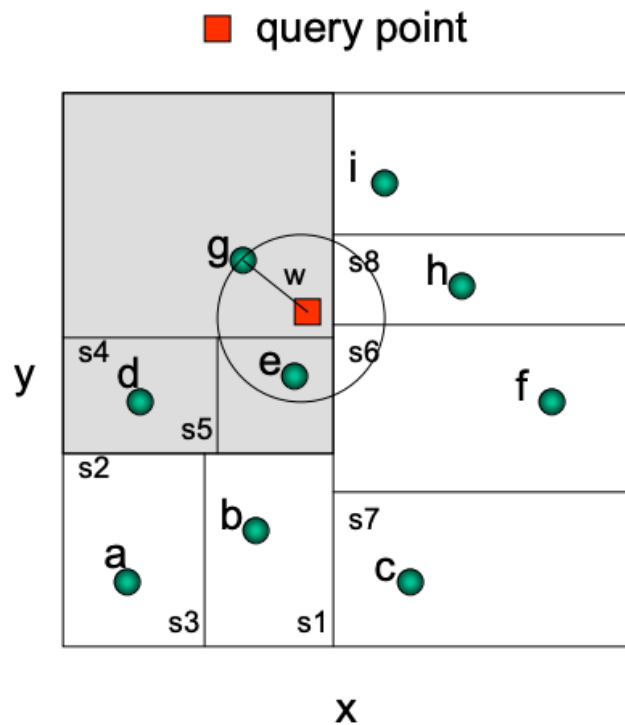


Kd-tree – kNN Search





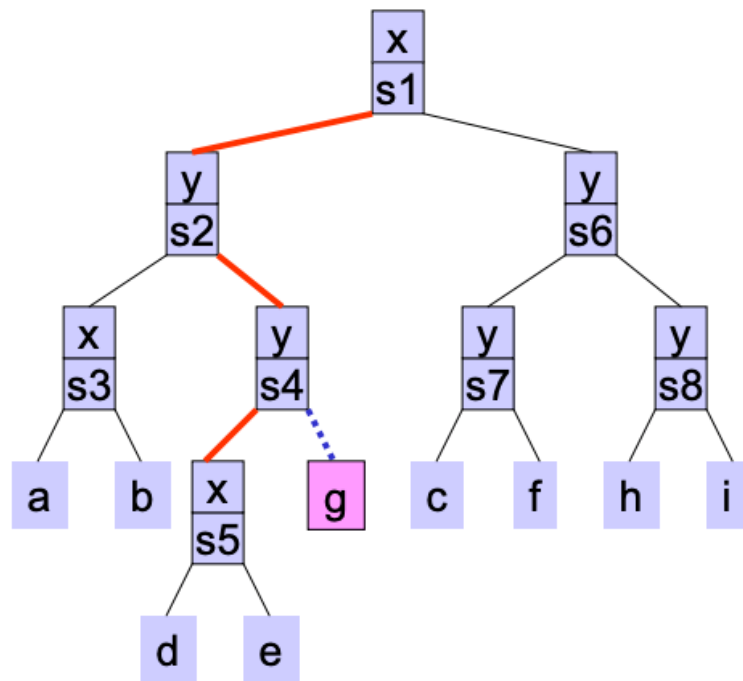
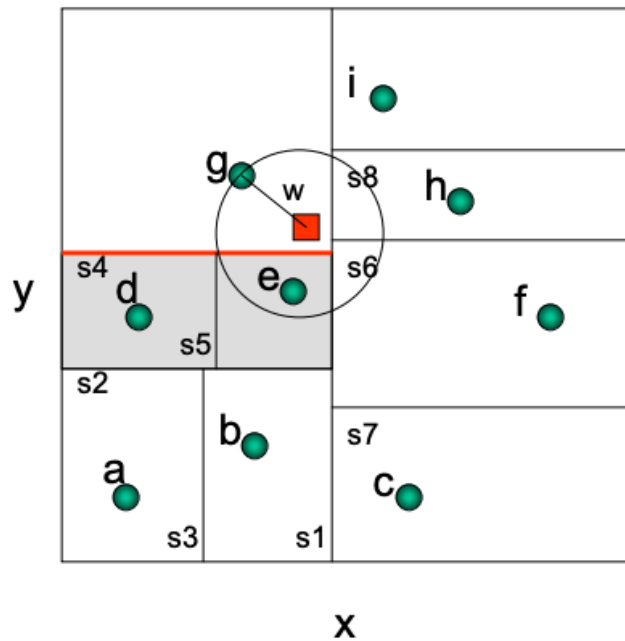
Kd-tree – kNN Search





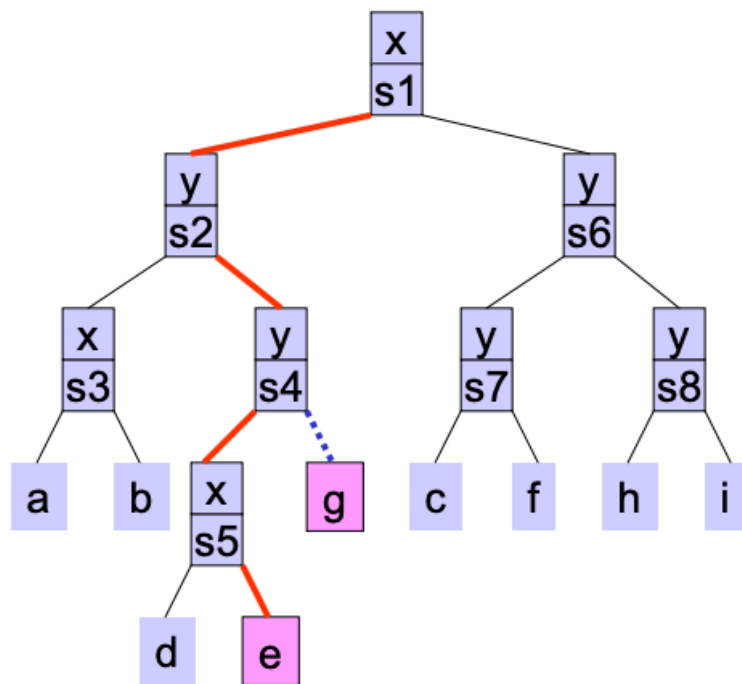
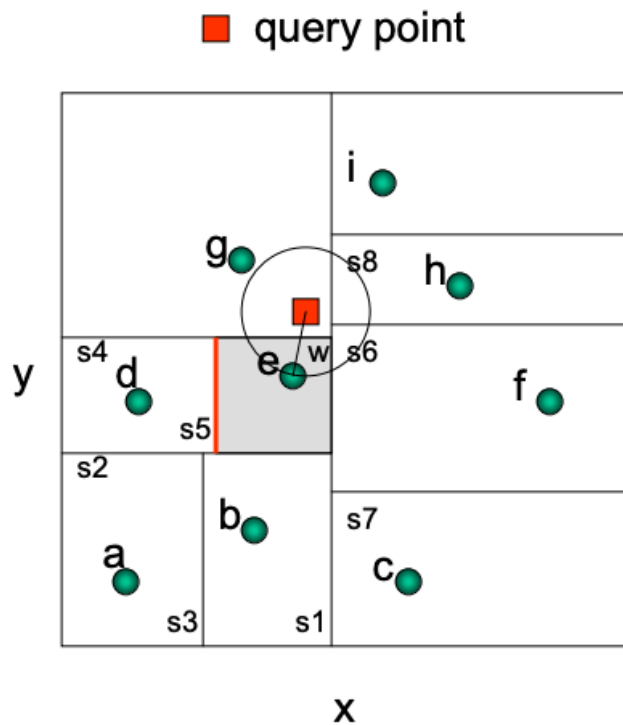
Kd-tree – kNN Search

■ query point



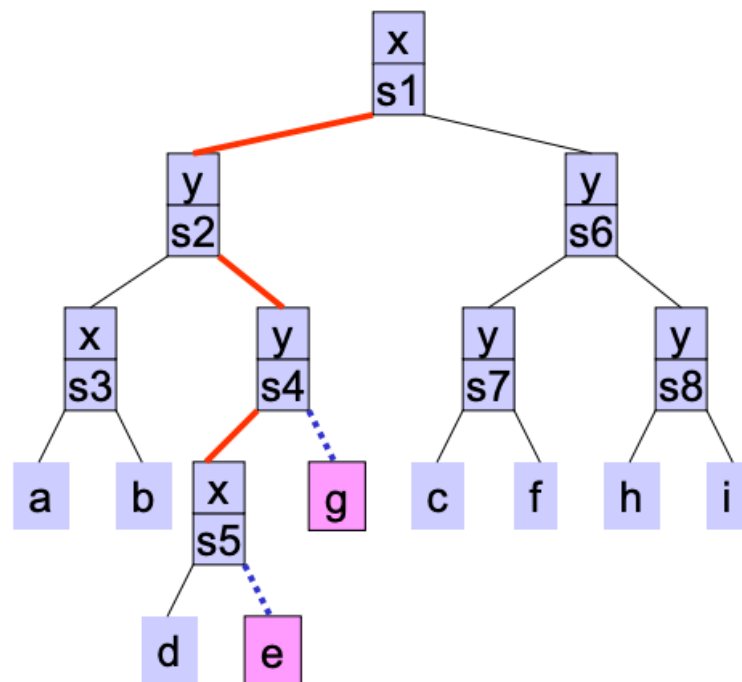
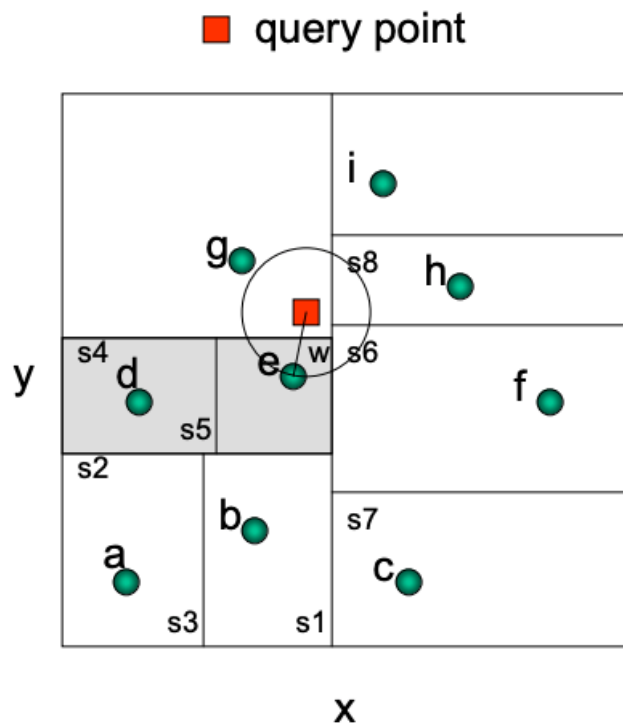


Kd-tree – kNN Search



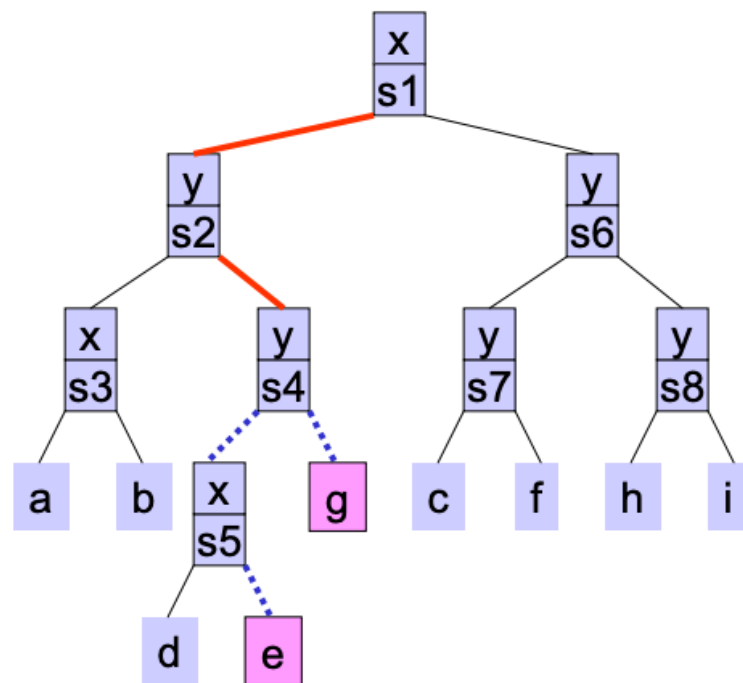
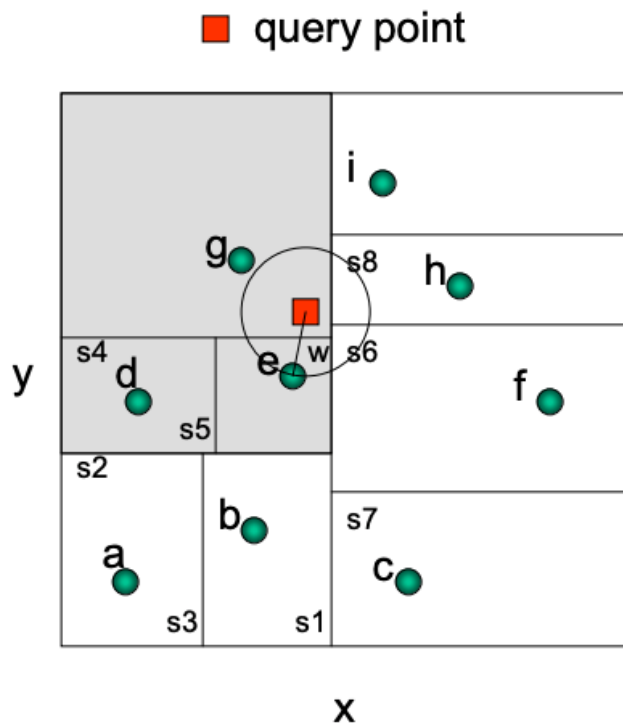


Kd-tree – kNN Search



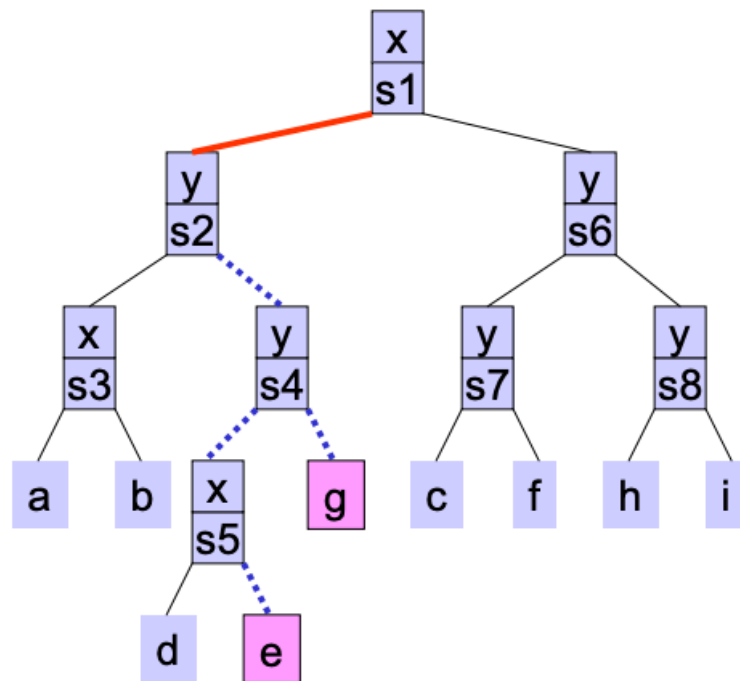
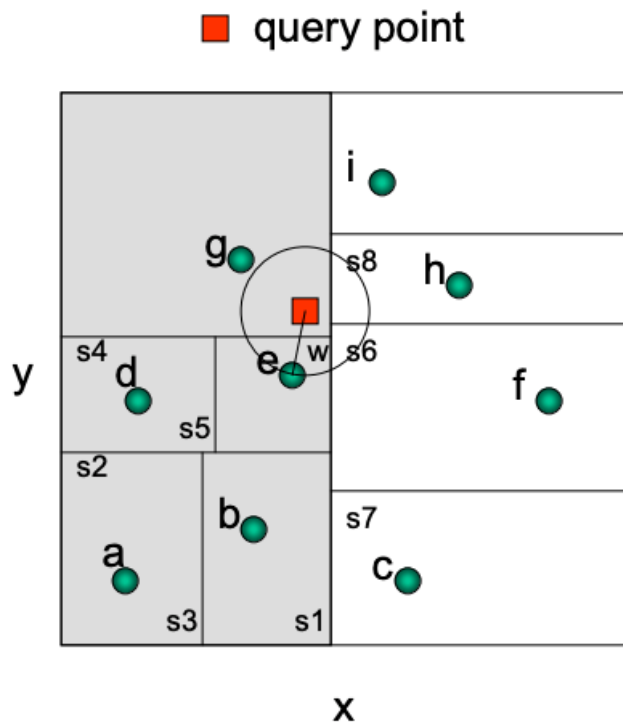


Kd-tree – kNN Search



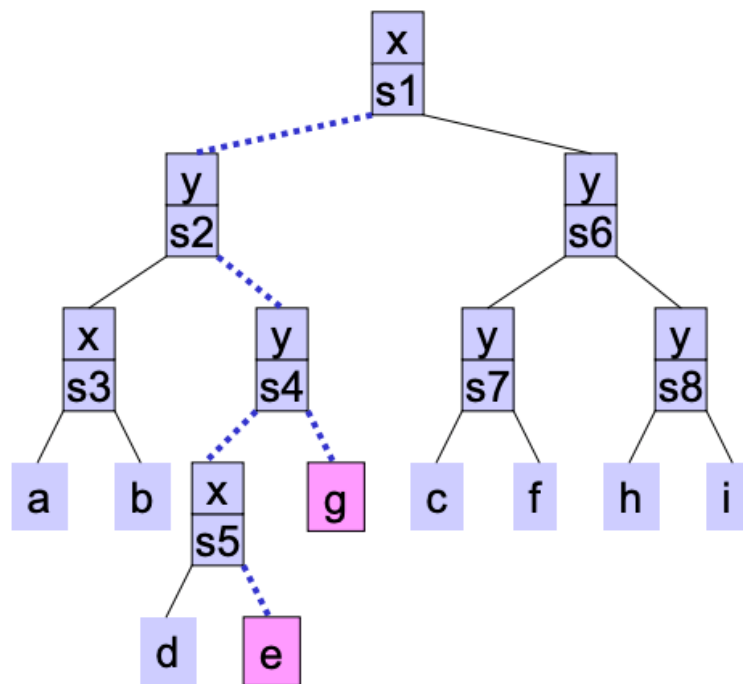
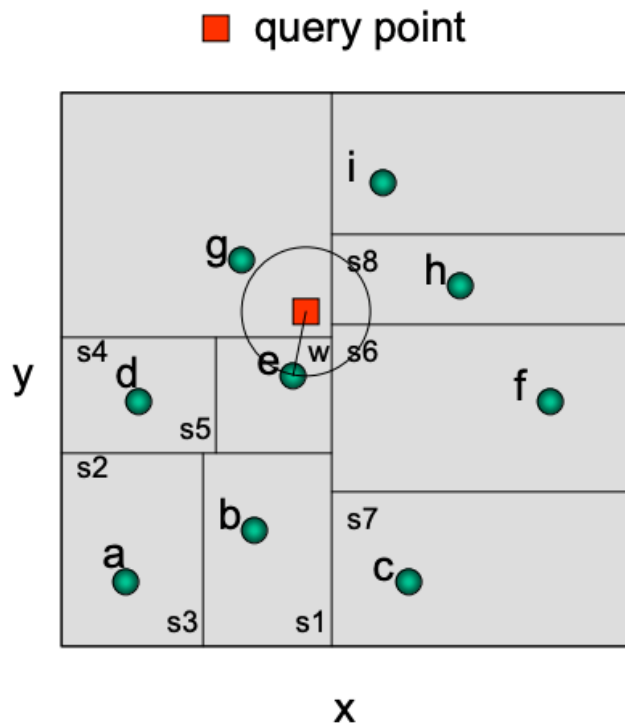


Kd-tree – kNN Search



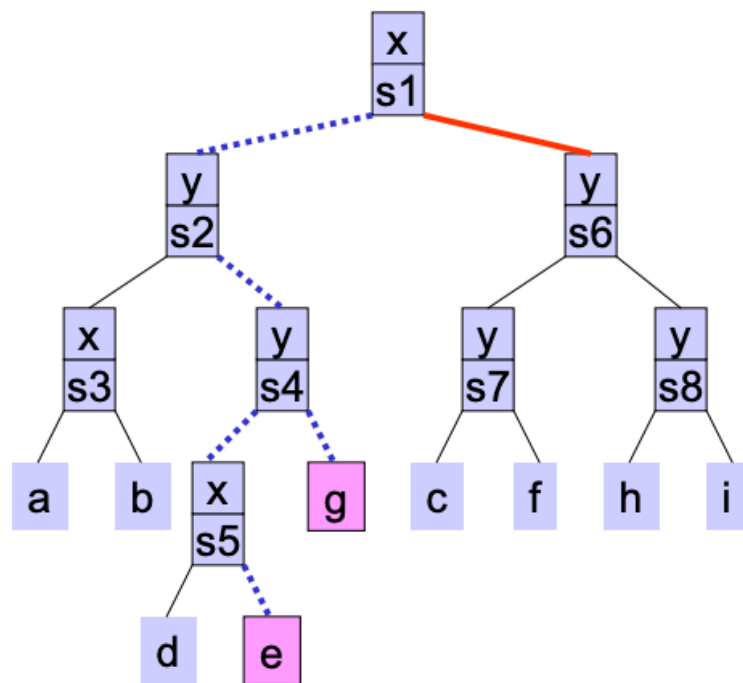
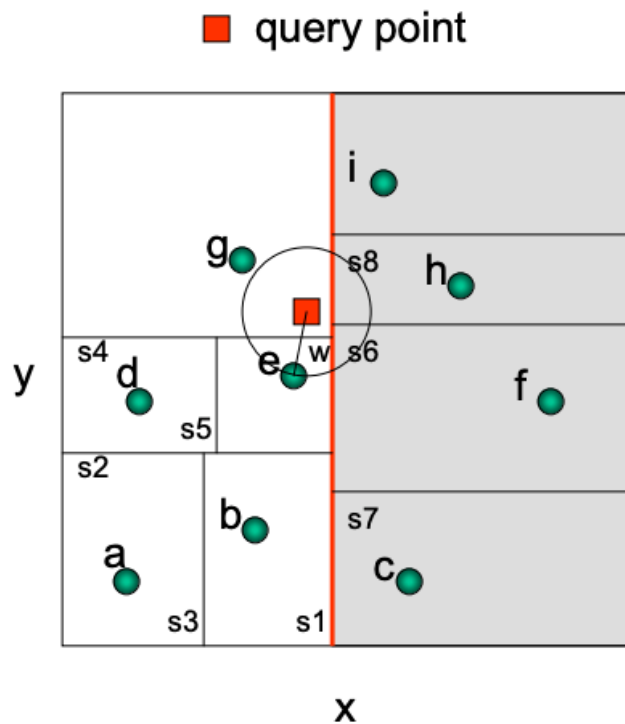


Kd-tree – kNN Search



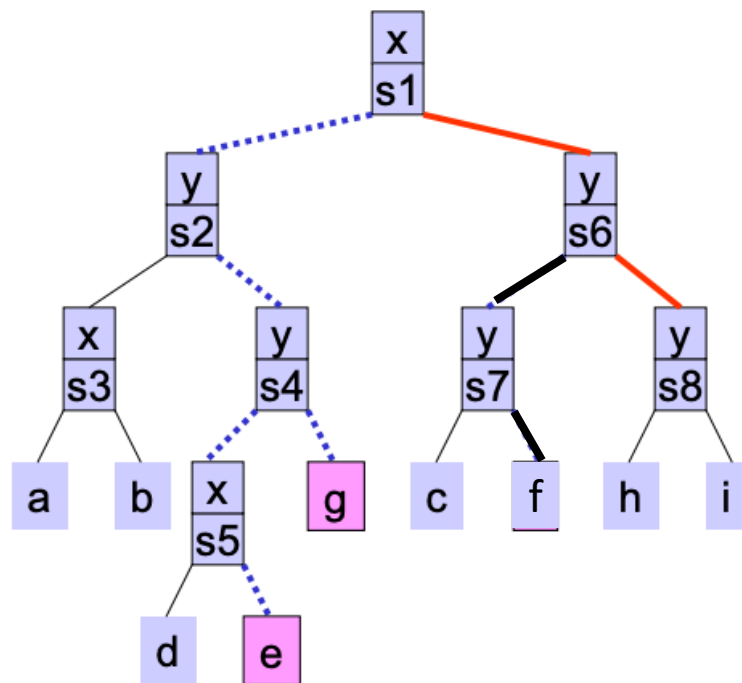
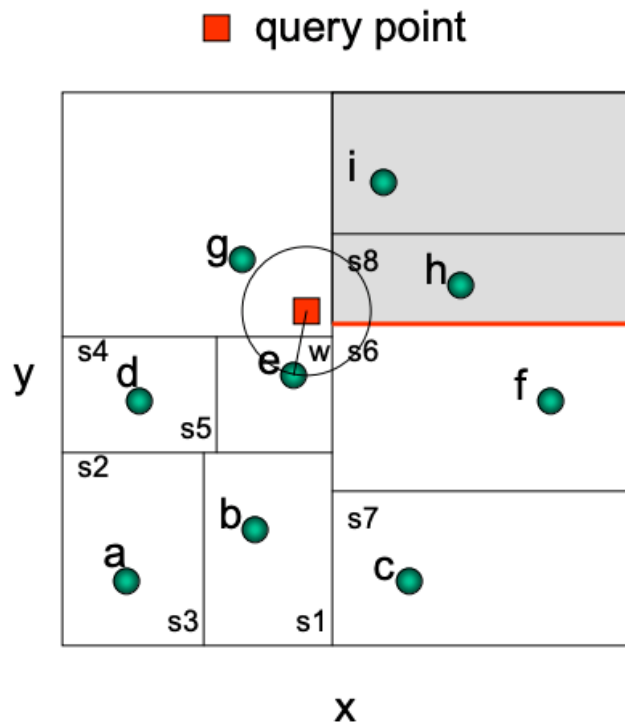


Kd-tree – kNN Search





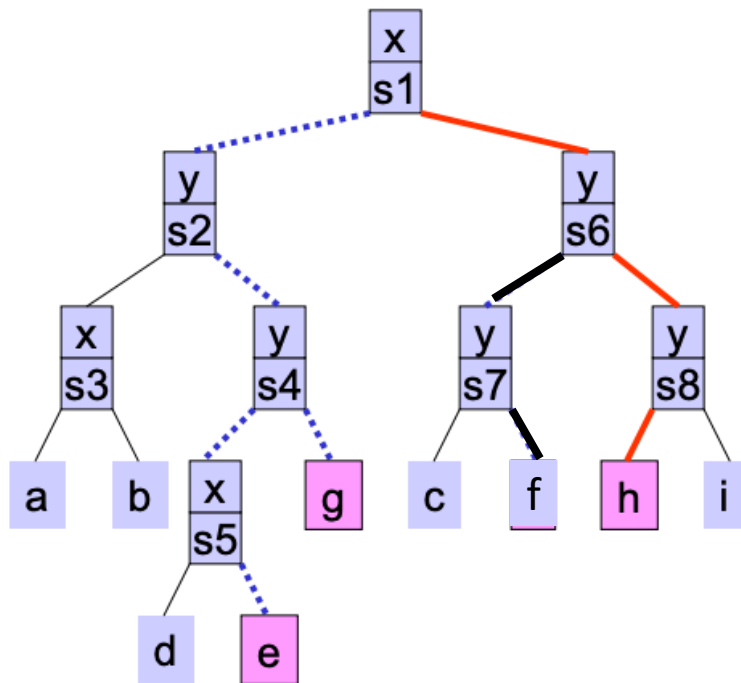
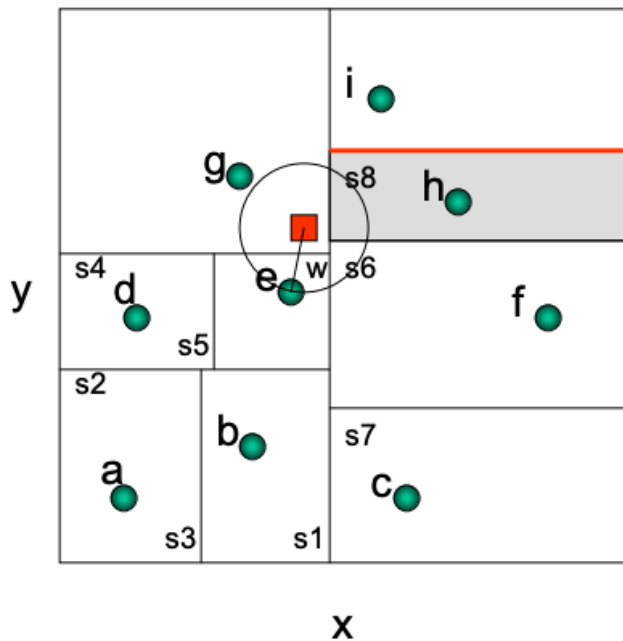
Kd-tree – kNN Search





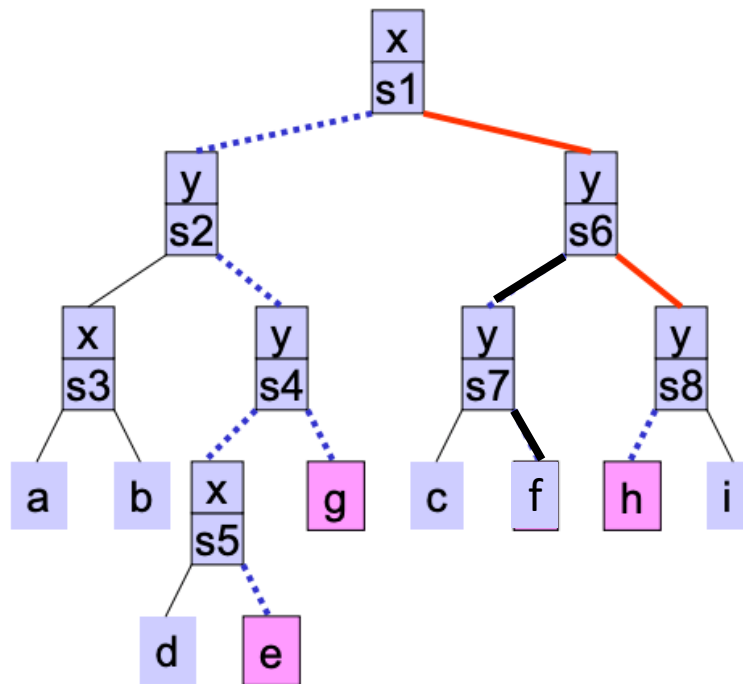
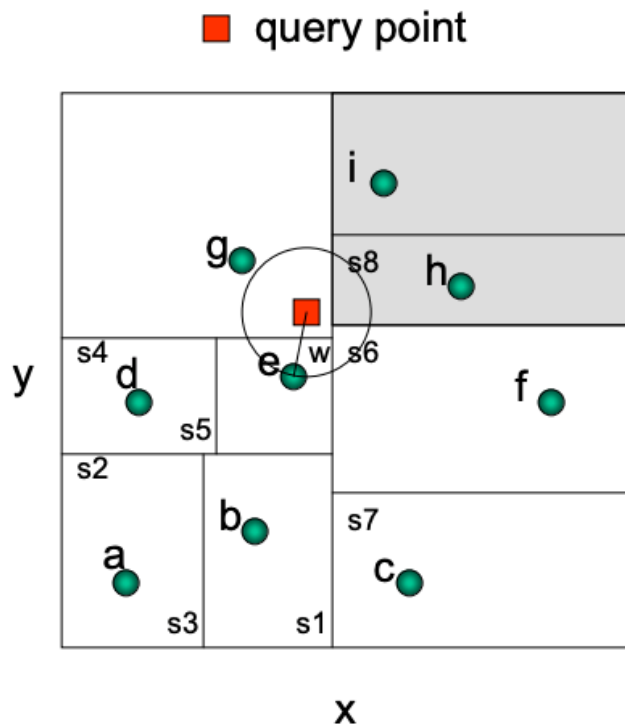
Kd-tree – kNN Search

■ query point



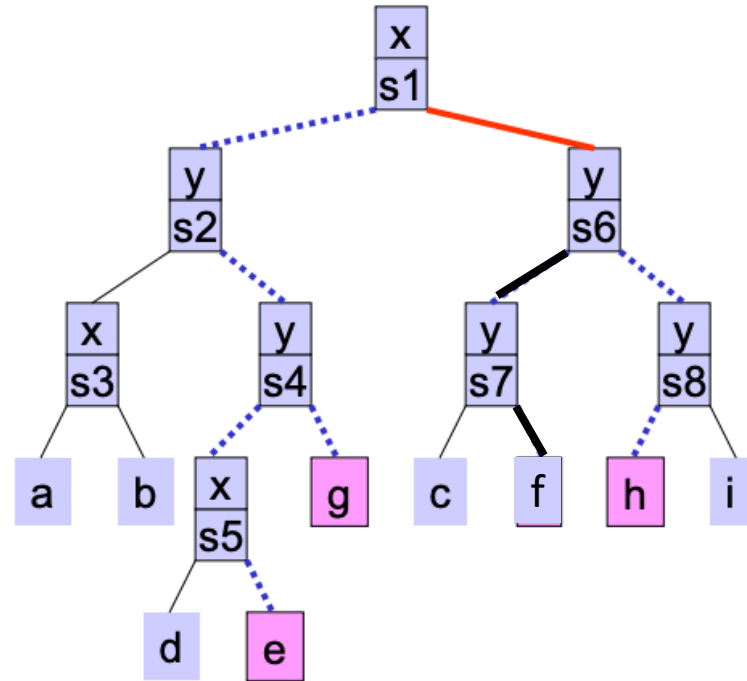
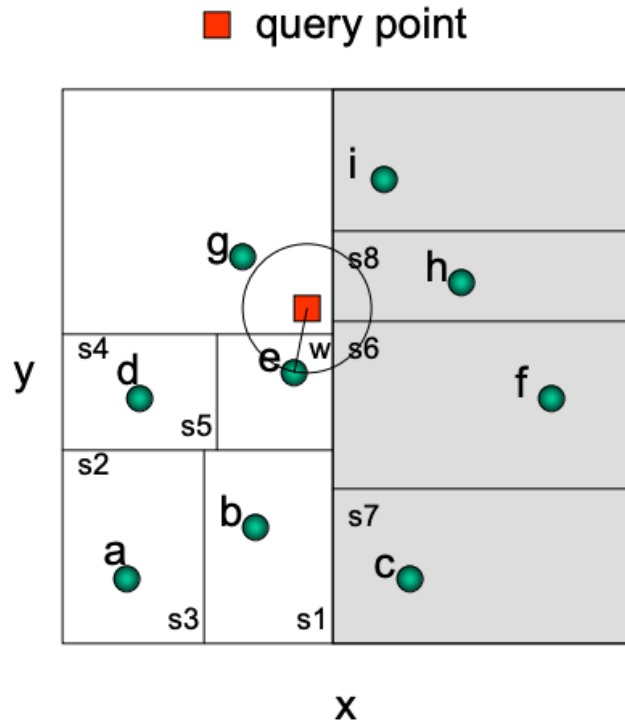


Kd-tree – kNN Search



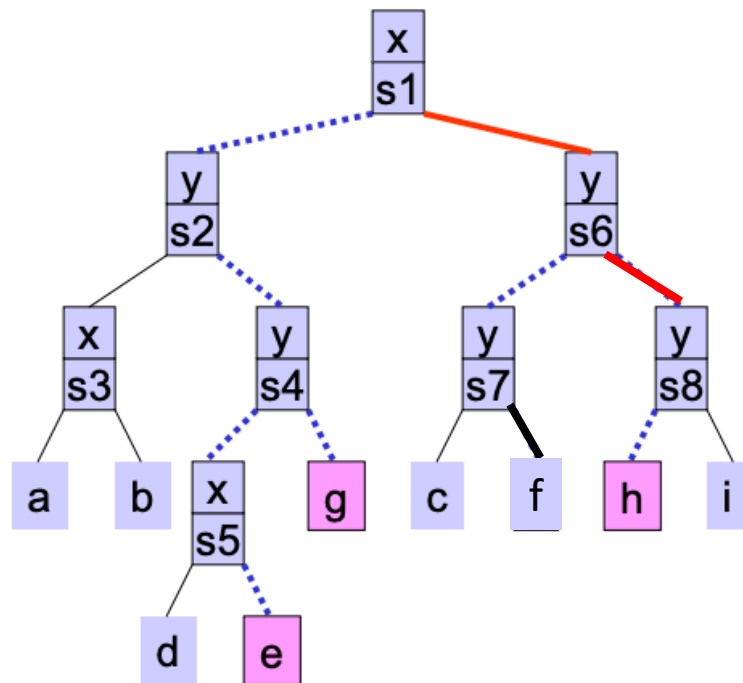
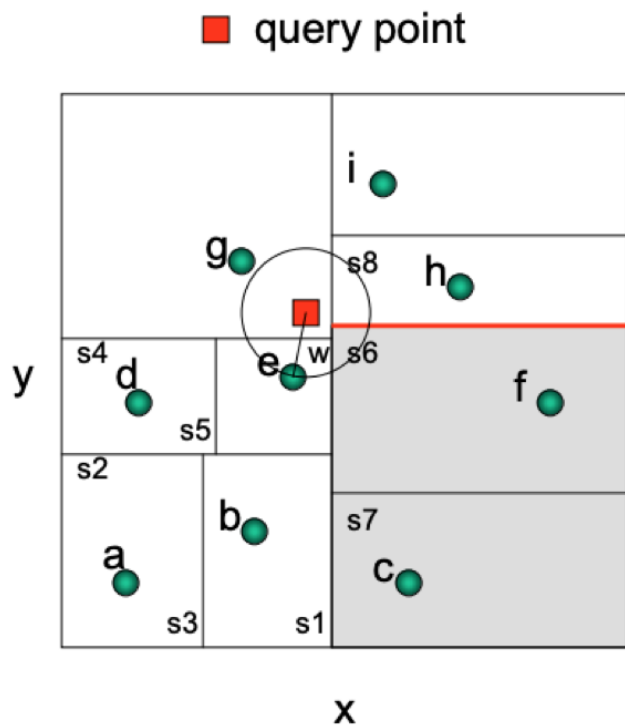


Kd-tree – kNN Search



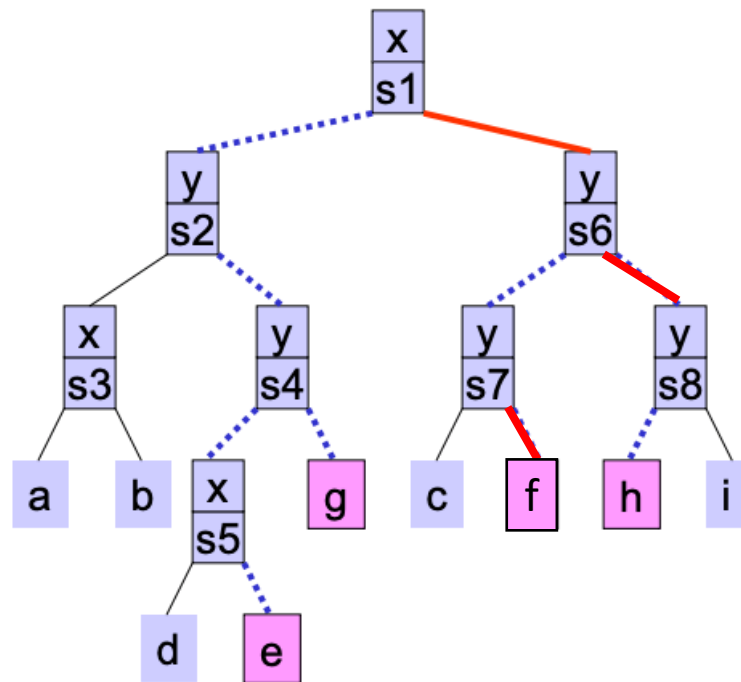
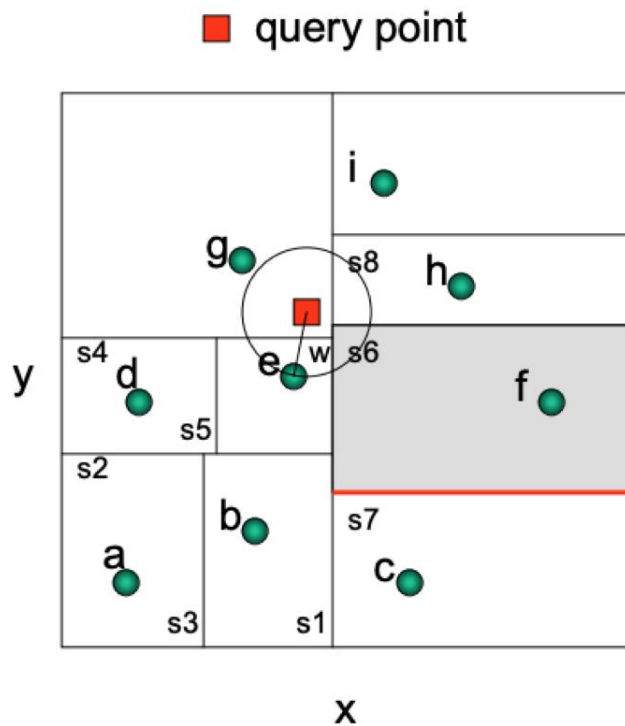


Kd-tree – kNN Search



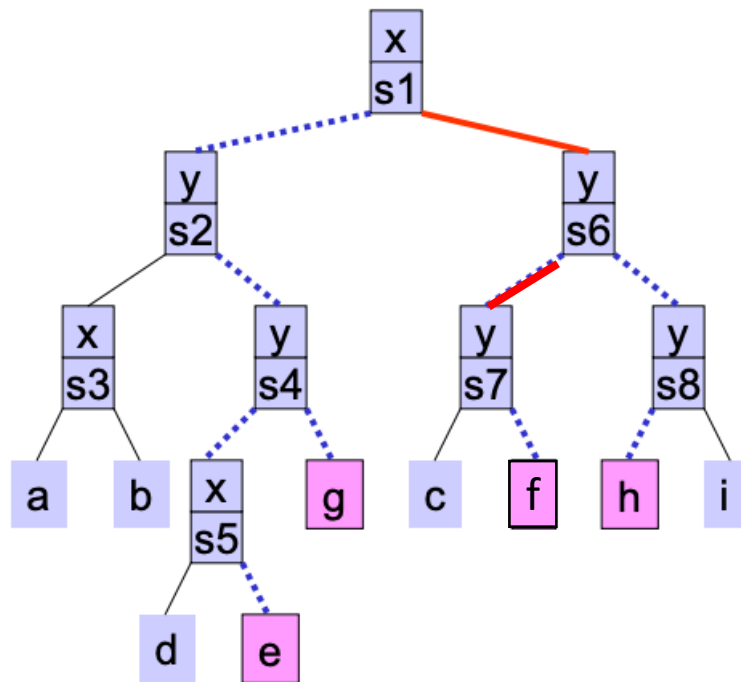
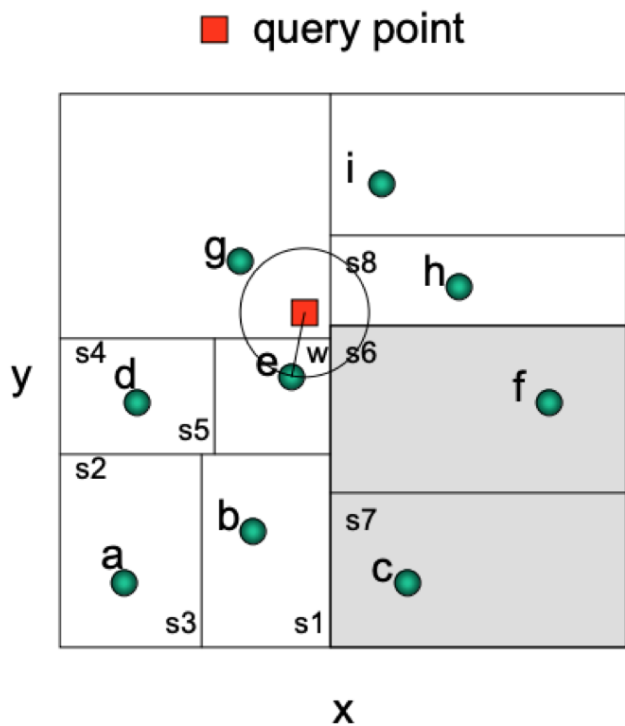


Kd-tree – kNN Search



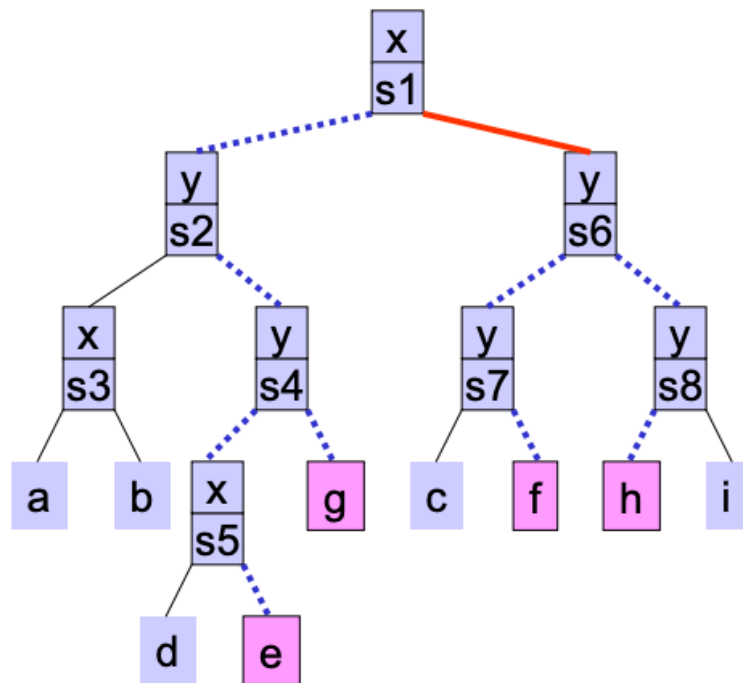
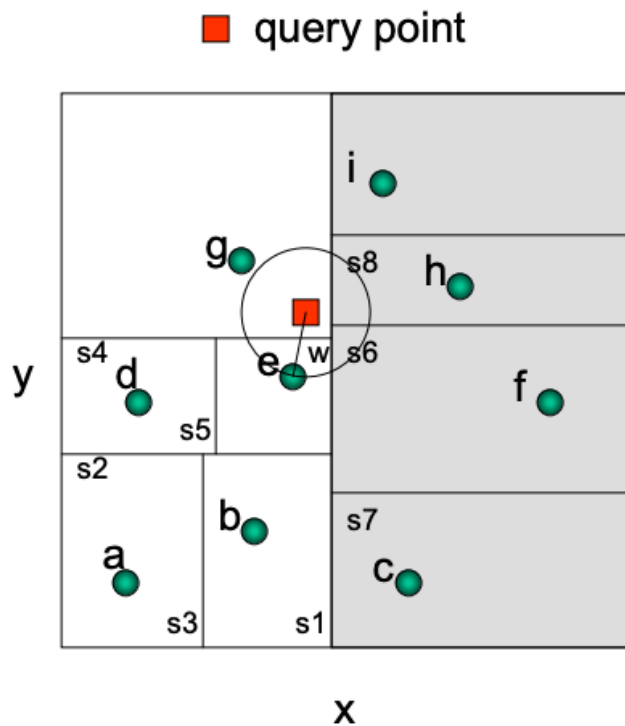


Kd-tree – kNN Search



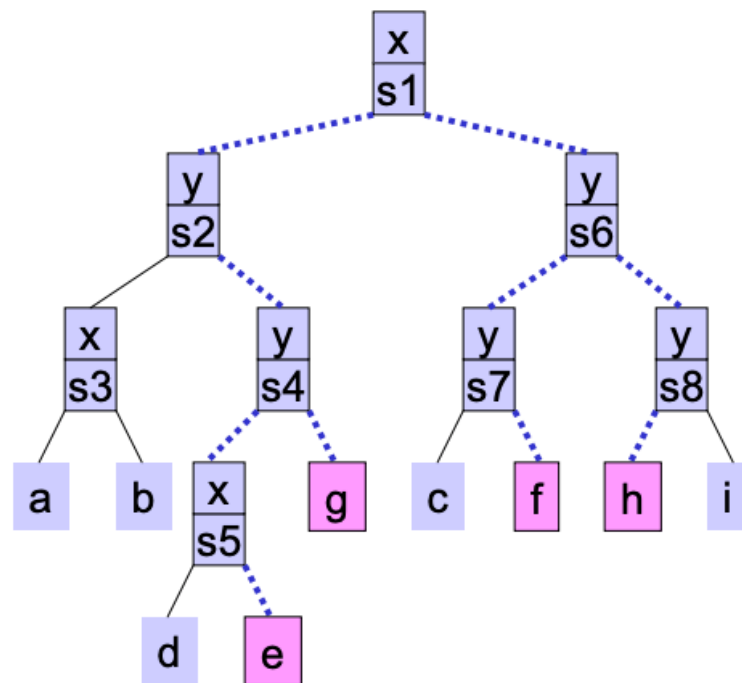
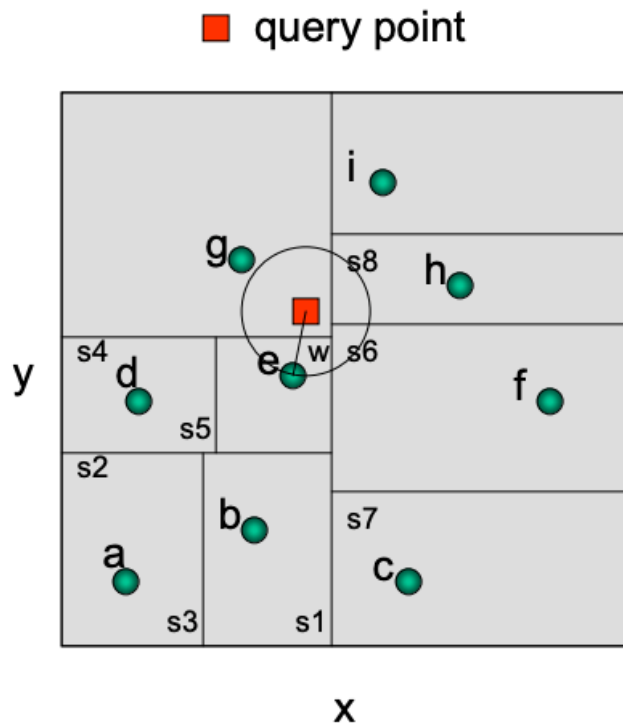


Kd-tree – kNN Search



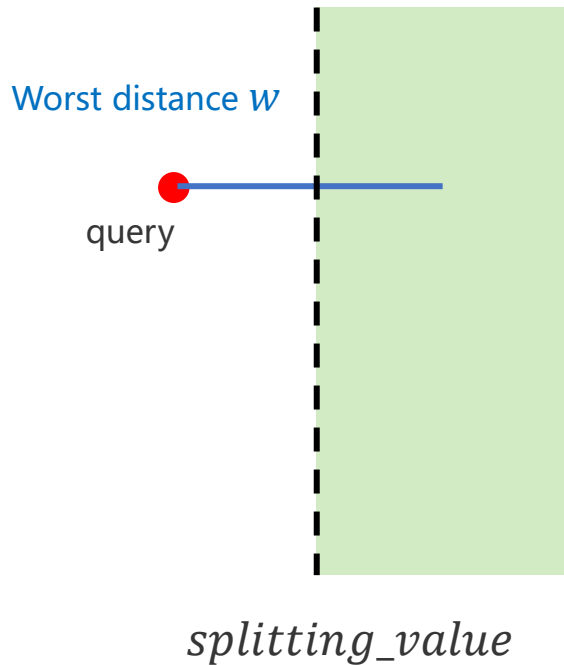


Kd-tree – kNN Search





Kd-tree – kNN Search



Criteria of a partition intersects with the worst-distance range:

$q[axis]$ inside the partition

OR

$$|q[axis] - splitting_value| < w$$

```
def knn_search(root: Node, db: np.ndarray, result_set: KNNResultSet, query: np.ndarray):  
    if root is None:  
        return False
```

```
    if root.is_leaf():
```

Compare query to every point inside the leaf, put into the result set

```
        # compare the contents of a leaf  
        leaf_points = db[root.point_indices, :]  
        diff = np.linalg.norm(np.expand_dims(query, 0) - leaf_points, axis=1)  
        for i in range(diff.shape[0]):  
            result_set.add_point(diff[i], root.point_indices[i])  
        return False
```

```
    if query[root.axis] <= root.value:
```

$q[\text{axis}]$ inside the partition

```
        knn_search(root.left, db, result_set, query)
```

```
        if math.fabs(query[root.axis] - root.value) < result_set.worstDist():  
            knn_search(root.right, db, result_set, query)
```

```
    else:
```

$|q[\text{axis}] - \text{splitting_value}| < w$

```
        knn_search(root.right, db, result_set, query)  
        if math.fabs(query[root.axis] - root.value) < result_set.worstDist():  
            knn_search(root.left, db, result_set, query)
```

```
    return False
```




Kd-tree Radius-NN Search



Exactly the same as kNN search except:

- Use *RadiusNNResultSet*, similar to BST search
- Fixed worst distance, instead of dynamic

```
1  if query[root.axis] <= root.value:
2      radius_search(root.left, db, result_set, query)
3      if math.fabs(query[root.axis] - root.value) < result_set.worstDist():
4          radius_search(root.right, db, result_set, query)
5  else:
6      radius_search(root.right, db, result_set, query)
7      if math.fabs(query[root.axis] - root.value) < result_set.worstDist():
8          radius_search(root.left, db, result_set, query)
9
10 return False
```



Kd-tree Search Complexity



1NN search is $O(\log n)$ for a balanced kd-tree



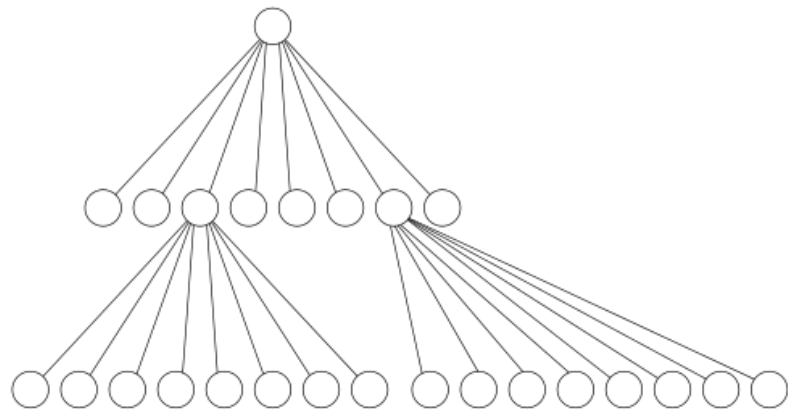
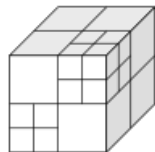
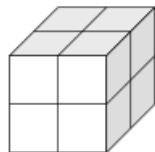
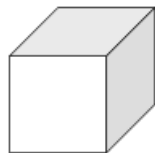
kNN/radiusNN complexity is hard to analyze

- Depends on the distribution of points
- Depends on k or r
- Varies from $O(\log n)$ to $O(n)$



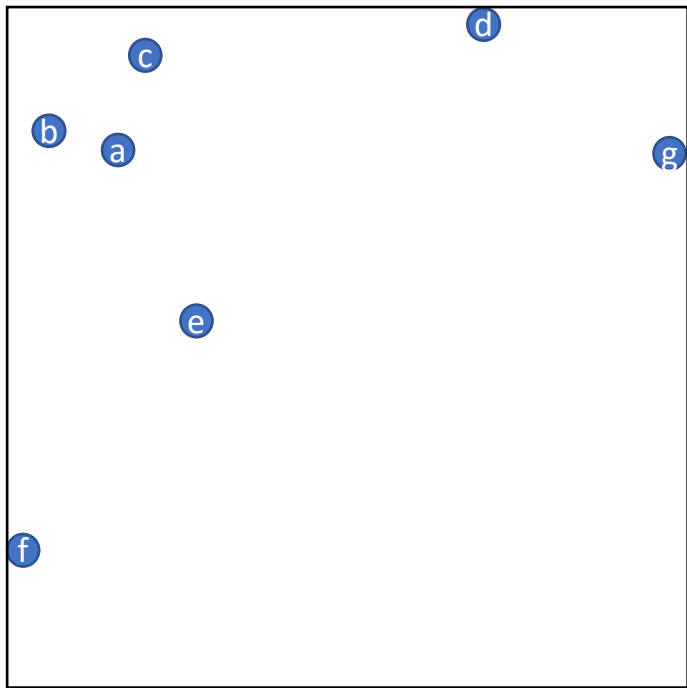
Octree

- Each node has 8 children
- oct – tree
- Specifically for 3D, $2^3=8$
- In kd-tree, it is non-trivial to determine whether the NN search is done, so we have to go back to root every time
- Octree is more efficient because we can stop without going back to root





Octree Construction

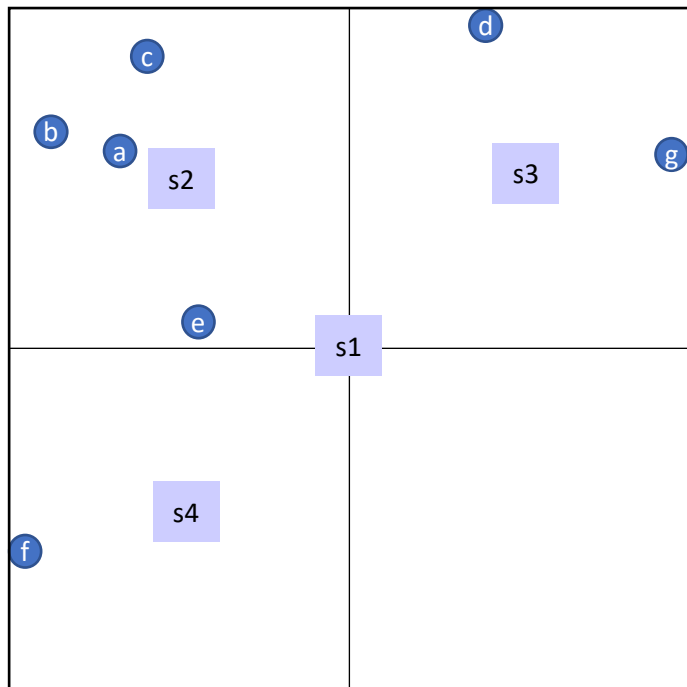


- Determine the extent of the first octant
- Octant is an element in the octree
- Octant is a cube.

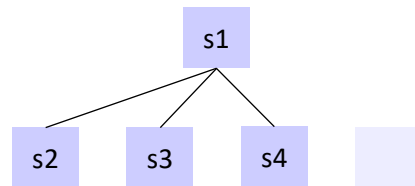
s1



Octree Construction

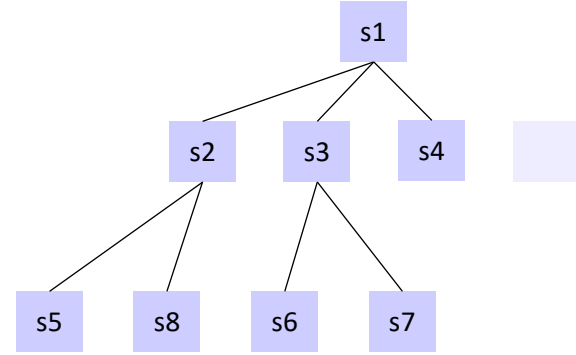
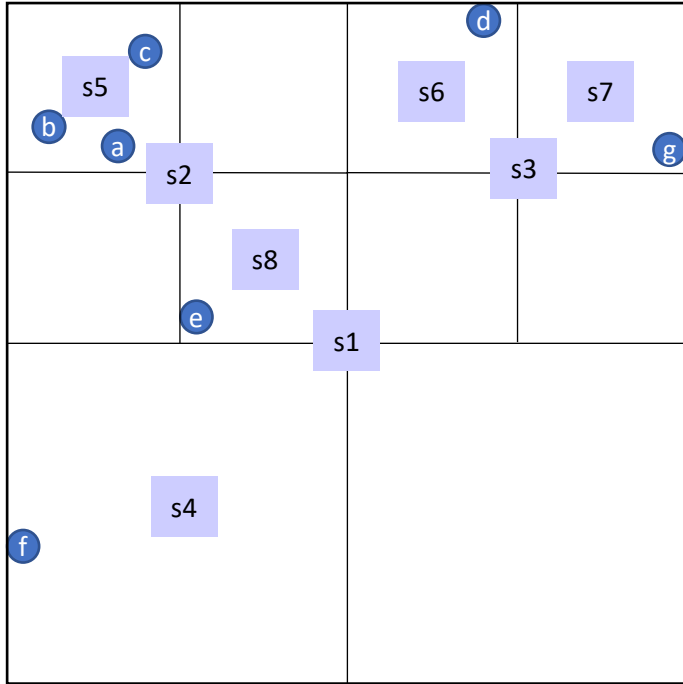


- Determine whether to further split the octant
- **leaf_size** = 1 here
- **min_extent** – avoid infinite splitting when there are repeated points



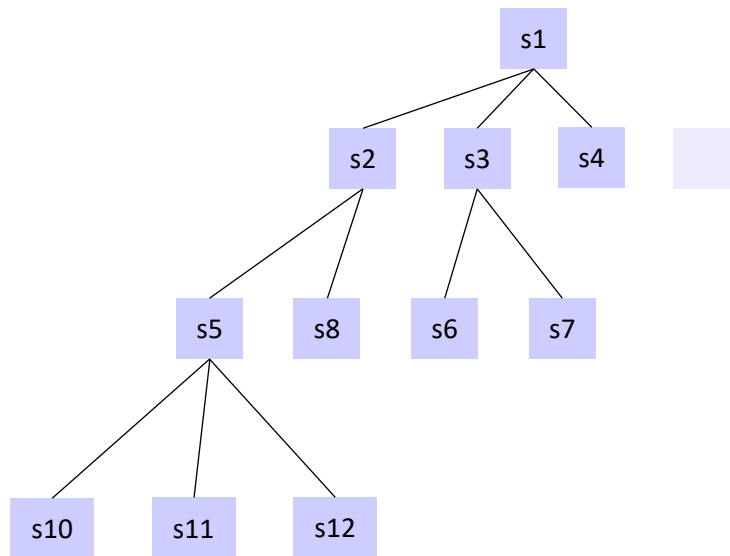
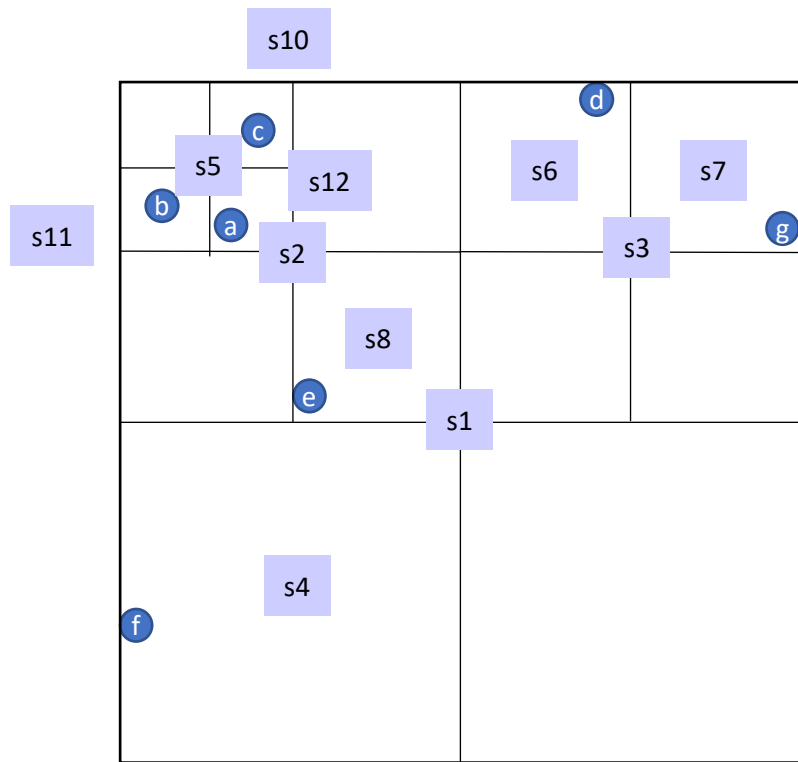


Octree Construction





Octree Construction





Octree Construction

Array of length 8 Center of the cube Point inside octant

```
class Octant:
    def __init__(self, children, center, extent, point_indices, is_leaf):
        self.children = children
        self.center = center
        self.extent = extent
        self.point_indices = point_indices
        self.is_leaf = is_leaf
```

0.5 * length


```

def octree_recursive_build(root, db, center, extent, point_indices, leaf_size, min_extent):
    if len(point_indices) == 0:
        return None

    if root is None:
        root = Octant([None for i in range(8)], center, extent, point_indices, is_leaf=True)

    # determine whether to split this octant
    if len(point_indices) <= leaf_size or extent <= min_extent:
        root.is_leaf = True
    else:
        root.is_leaf = False
        children_point_indices = [[] for i in range(8)]
        for point_idx in point_indices:
            point_db = db[point_idx]
            morton_code = 0
            if point_db[0] > center[0]:
                morton_code = morton_code | 1
            if point_db[1] > center[1]:
                morton_code = morton_code | 2
            if point_db[2] > center[2]:
                morton_code = morton_code | 4
            children_point_indices[morton_code].append(point_idx)

        # create children
        factor = [-0.5, 0.5]
        for i in range(8):
            child_center_x = center[0] + factor[(i & 1) > 0] * extent
            child_center_y = center[1] + factor[(i & 2) > 0] * extent
            child_center_z = center[2] + factor[(i & 4) > 0] * extent
            child_extent = 0.5 * extent
            child_center = np.asarray([child_center_x, child_center_y, child_center_z])
            root.children[i] = octree_recursive_build(root.children[i],
                                                    db,
                                                    child_center,
                                                    child_extent,
                                                    children_point_indices[i],
                                                    leaf_size,
                                                    min_extent)

    return root

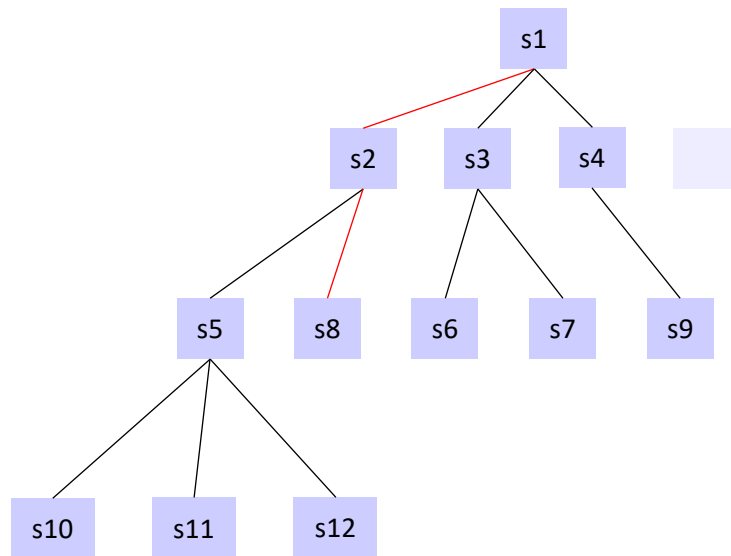
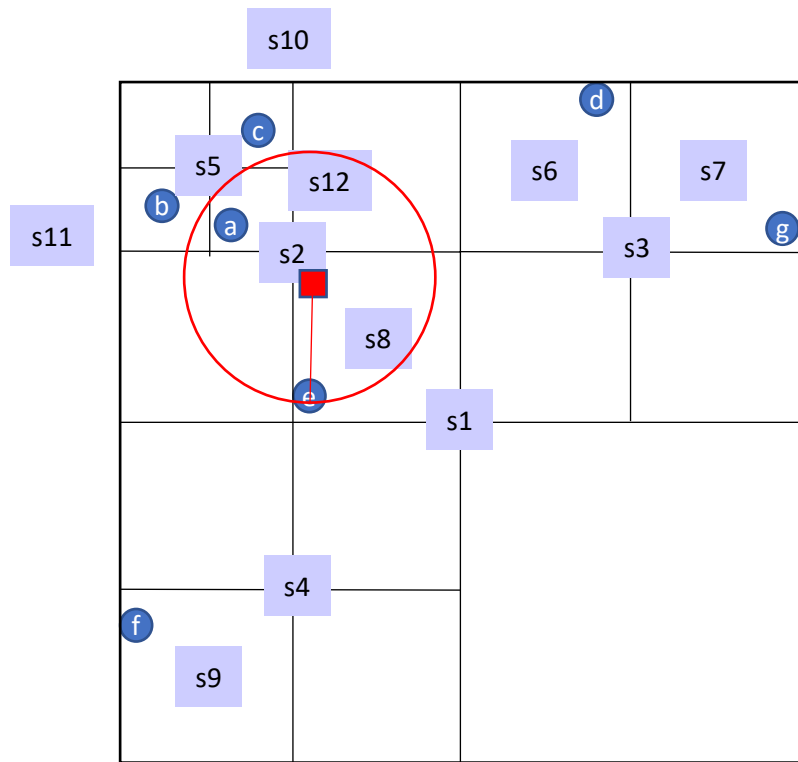
```

Determine which child a point belongs to

Determine child center & extent

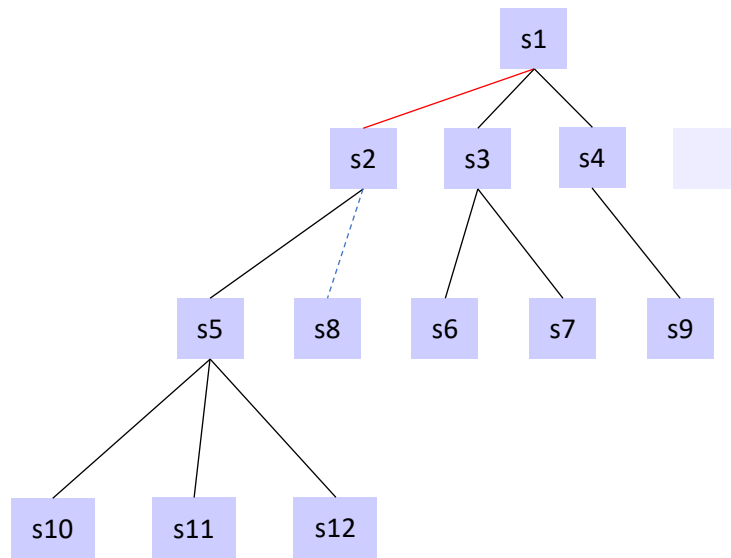
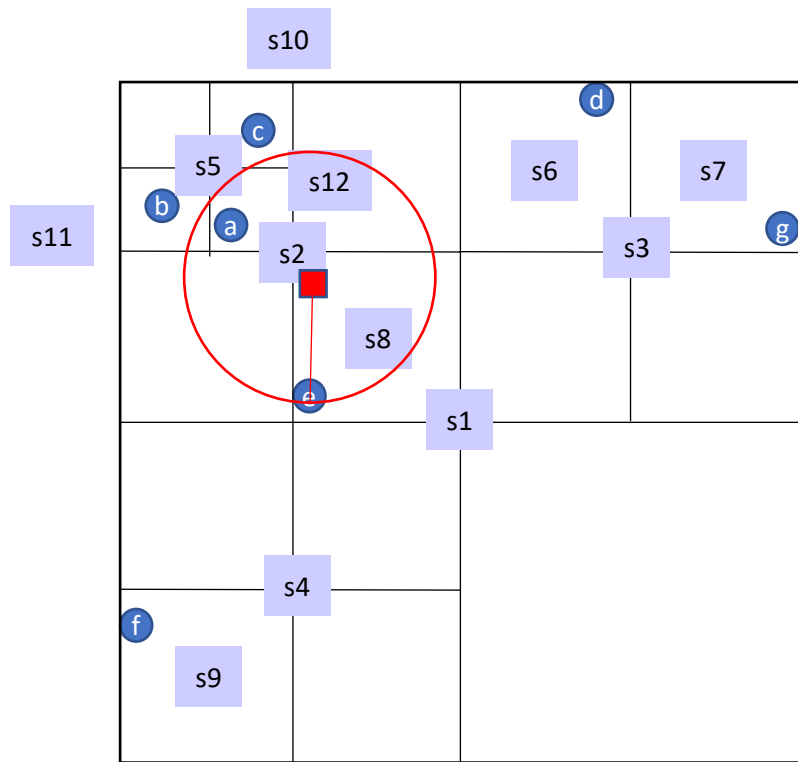


Octree kNN Search



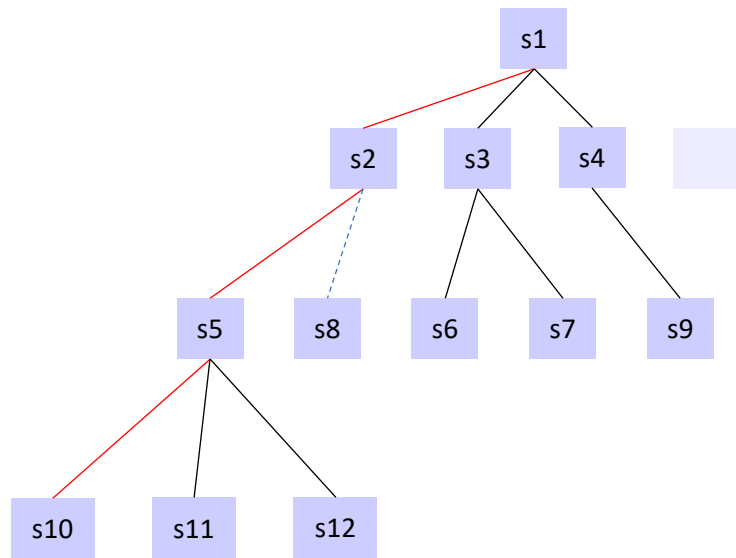
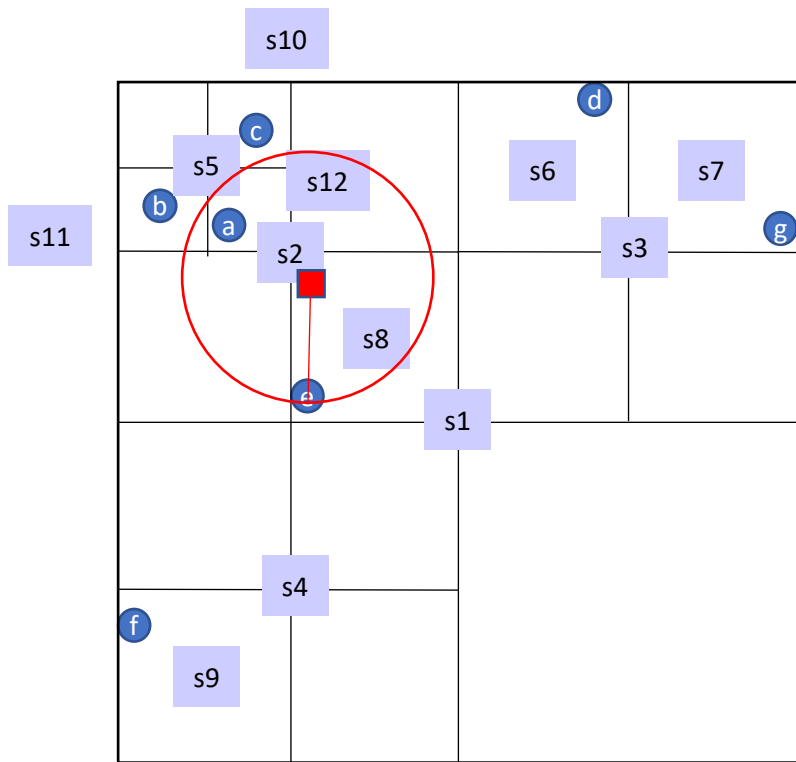


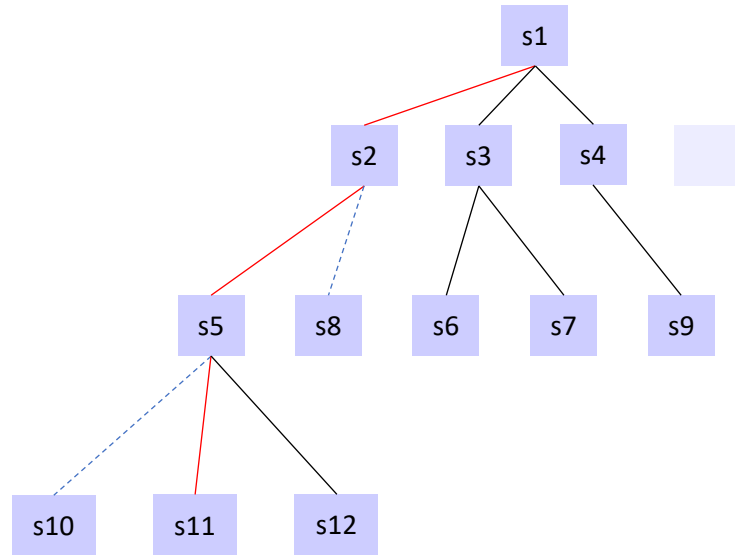
Octree kNN Search





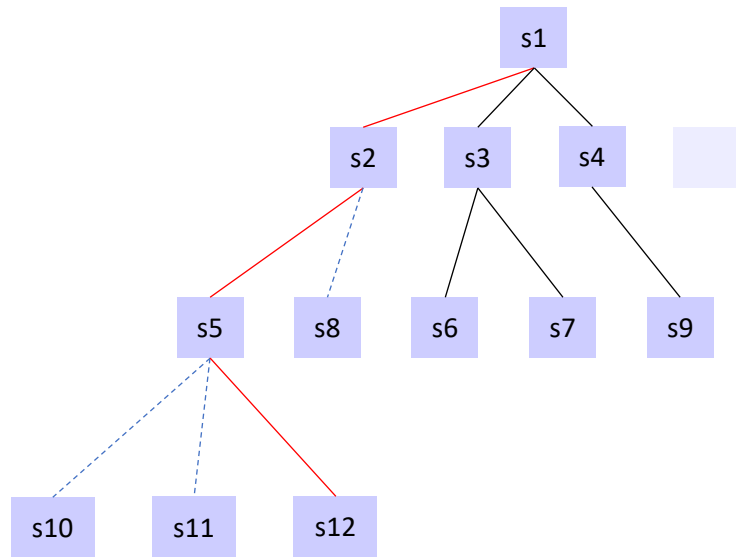
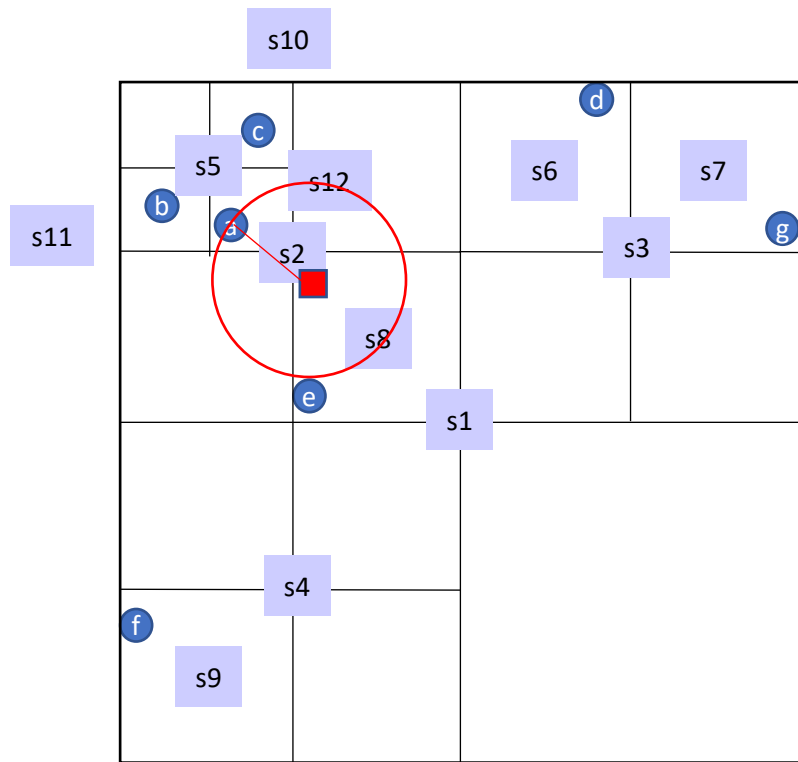
Octree kNN Search





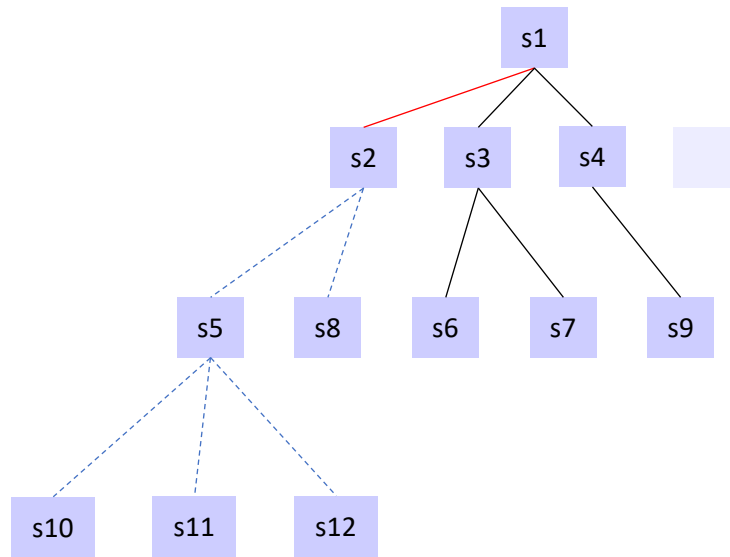
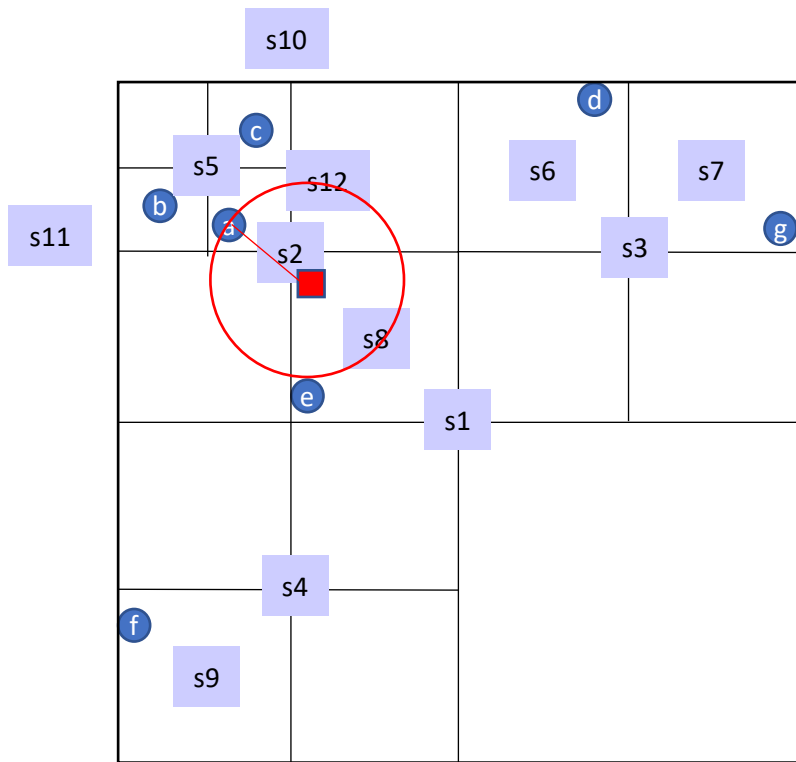


Octree kNN Search



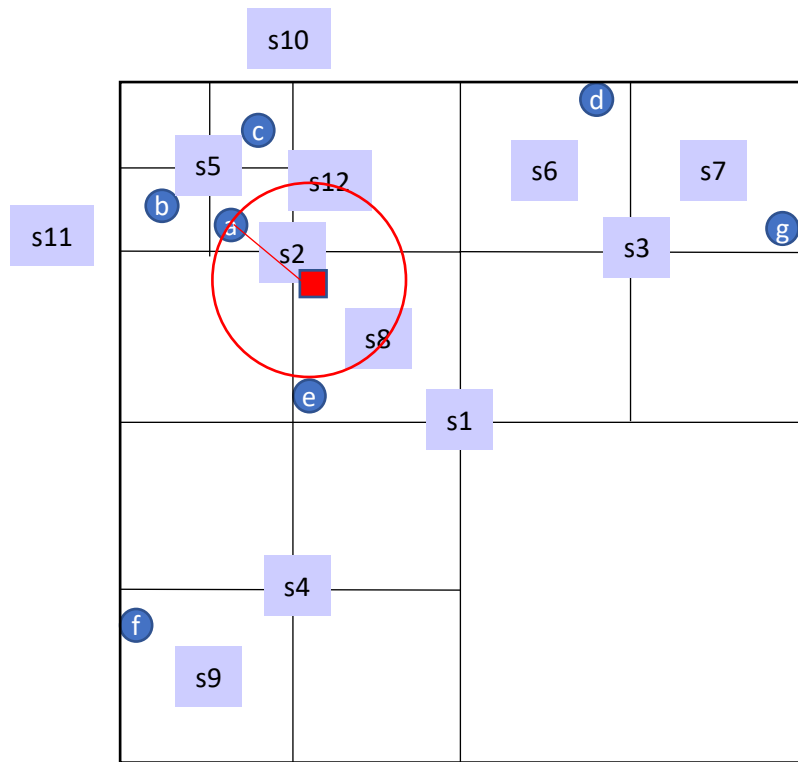


Octree kNN Search

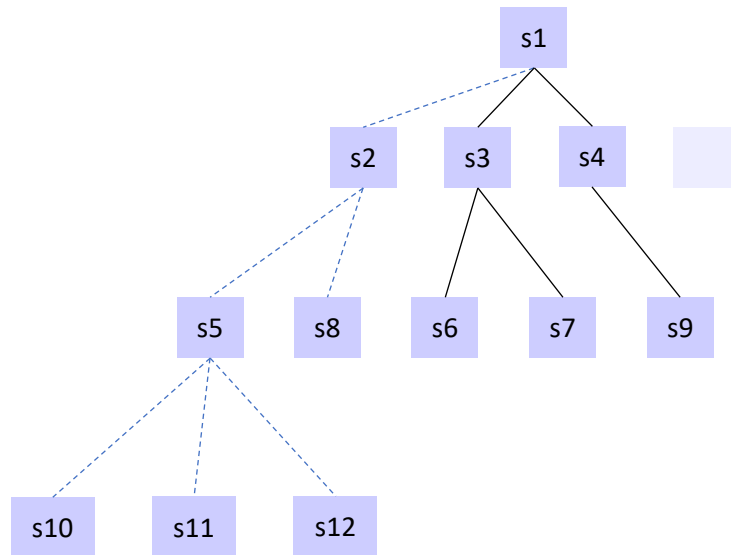




Octree kNN Search



Query ball inside s2, search end!




```
def octree_knn_search(root: Octant, db: np.ndarray, result_set: KNNResultSet, query: np.ndarray):
```

```
    if root is None:  
        return False
```

```
    if root.is_leaf and len(root.point_indices) > 0:
```

```
        # compare the contents of a leaf  
        leaf_points = db[root.point_indices, :]  
        diff = np.linalg.norm(np.expand_dims(query, 0) - leaf_points, axis=1)  
        for i in range(diff.shape[0]):  
            result_set.add_point(diff[i], root.point_indices[i])  
        # check whether we can stop search now  
        return inside(query, result_set.worstDist(), root)
```

Compare all points in a leaf

```
    # go to the relevant child first
```

```
    morton_code = 0
```

```
    if query[0] > root.center[0]:  
        morton_code = morton_code | 1
```

```
    if query[1] > root.center[1]:  
        morton_code = morton_code | 2
```

```
    if query[2] > root.center[2]:  
        morton_code = morton_code | 4
```

Determine & search the most relevant child

```
    if octree_knn_search(root.children[morton_code], db, result_set, query):  
        return True
```

```
    # check other children
```

```
    for c, child in enumerate(root.children):  
        if c == morton_code or child is None:  
            continue
```

If an octant is not overlapping with query ball, skip

```
        if False == overlaps(query, result_set.worstDist(), child):  
            continue
```

```
        if octree_knn_search(child, db, result_set, query):  
            return True
```

If query ball is inside an octant, stop

```
    # final check of if we can stop search
```

```
    return inside(query, result_set.worstDist(), root)
```



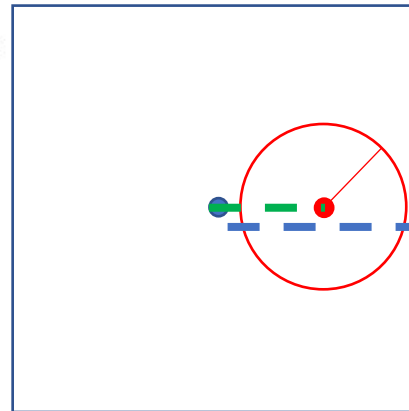
Function *inside*

```
def inside(query: np.ndarray, radius: float, octant: Octant):  
    """  
    Determines if the query ball is inside the octant  
    :param query:  
    :param radius:  
    :param octant:  
    :return:  
    """  
    query_offset = query - octant.center  
    query_offset_abs = np.fabs(query_offset)  
    possible_space = query_offset_abs + radius  
    return np.all(possible_space < octant.extent)
```

Green dash line

Blue dash line

Red line





Function *overlaps*

```
def overlaps(query: np.ndarray, radius: float, octant: Octant):
```

```
    """
```

```
    Determines if the query ball overlaps with the octant
```

```
    :param query:
```

```
    :param radius:
```

```
    :param octant:
```

```
    :return:
```

```
    """
```

```
    query_offset = query - octant.center
```

```
    query_offset_abs = np.fabs(query_offset)
```

```
    # completely outside, since query is outside the relevant area
```

```
    max_dist = radius + octant.extent
```

```
    if np.any(query_offset_abs > max_dist):
```

```
        return False
```

Case 1

```
    # if pass the above check, consider the case that the ball is contacting the face of the octant
```

```
    if np.sum((query_offset_abs < octant.extent).astype(np.int)) >= 2:
```

```
        return True
```

Case 2

```
    # consider the case that the ball is contacting the edge or corner of the octant
```

```
    # since the case of the ball center (query) inside octant has been considered,
```

```
    # we only consider the ball center (query) outside octant
```

```
    x_diff = max(query_offset_abs[0] - octant.extent, 0)
```

```
    y_diff = max(query_offset_abs[1] - octant.extent, 0)
```

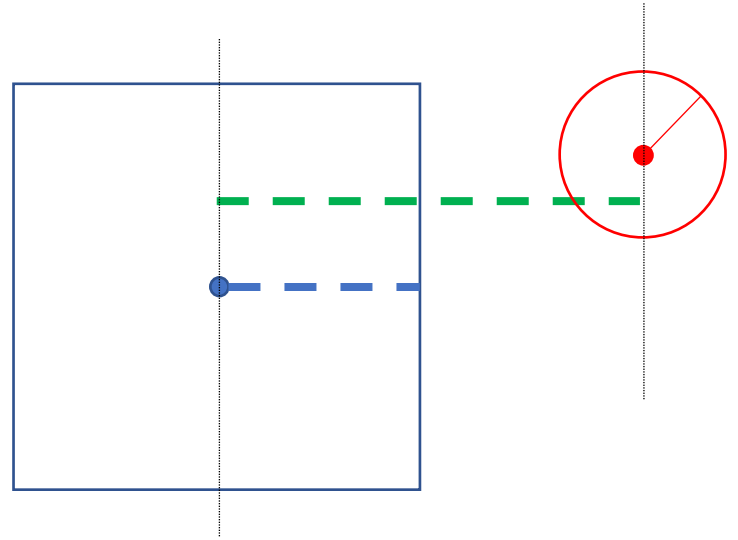
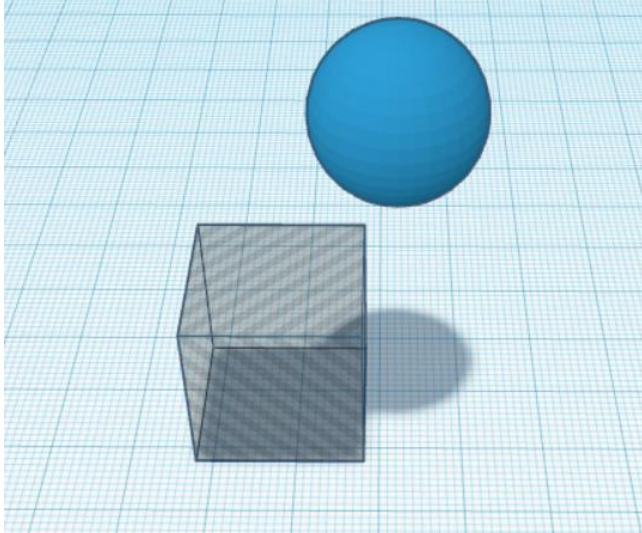
```
    z_diff = max(query_offset_abs[2] - octant.extent, 0)
```

Case 3

```
    return x_diff * x_diff + y_diff * y_diff + z_diff * z_diff < radius * radius
```



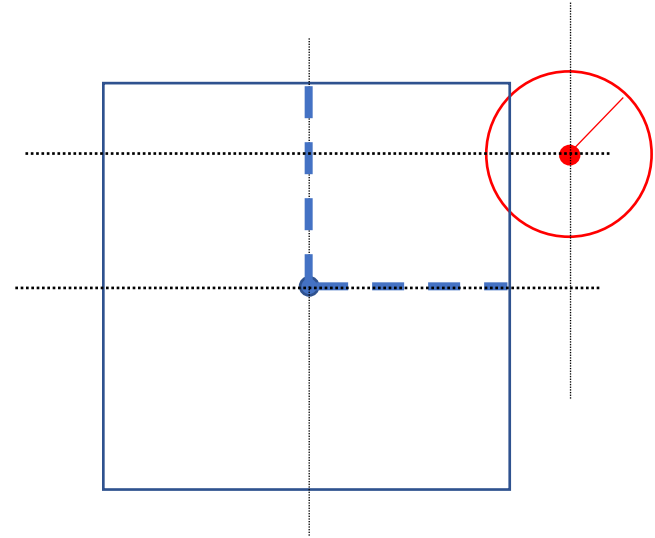
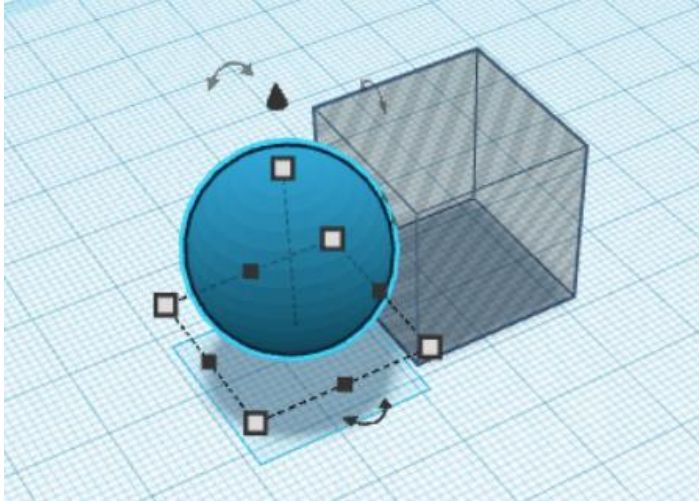
Function *overlaps* – Case 1



```
# completely outside, since query is outside the relevant area
max_dist = radius + octant.extent
if np.any(query_offset_abs > max_dist):
    return False
```



Function *overlaps* – Case 2

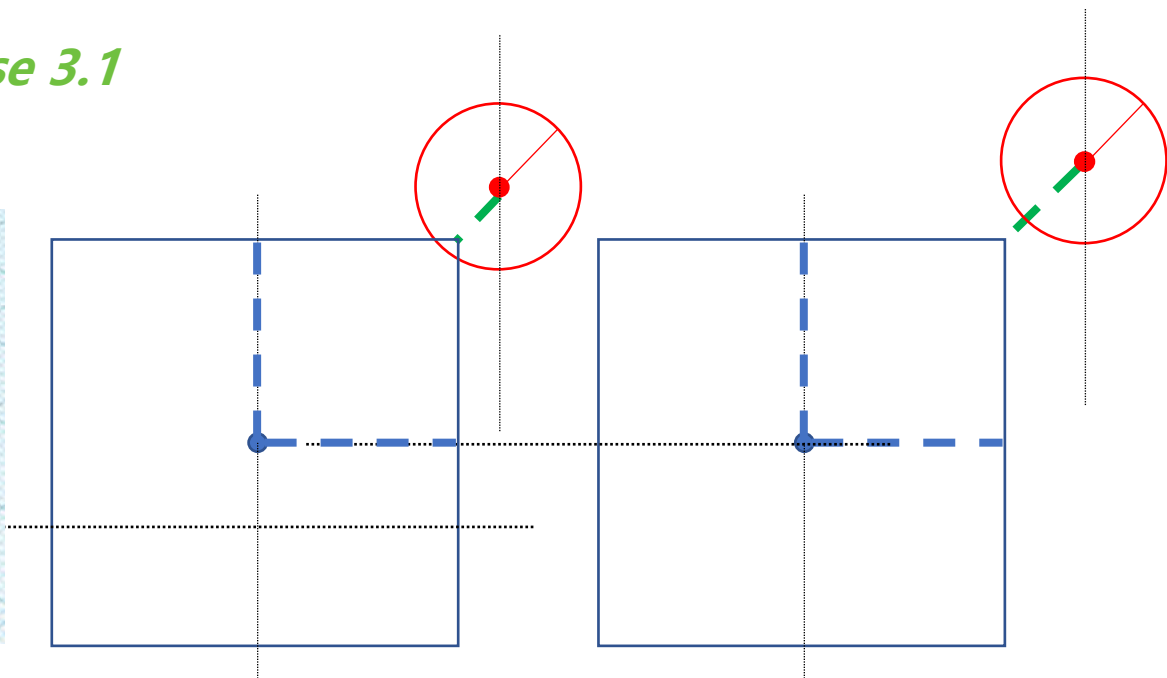
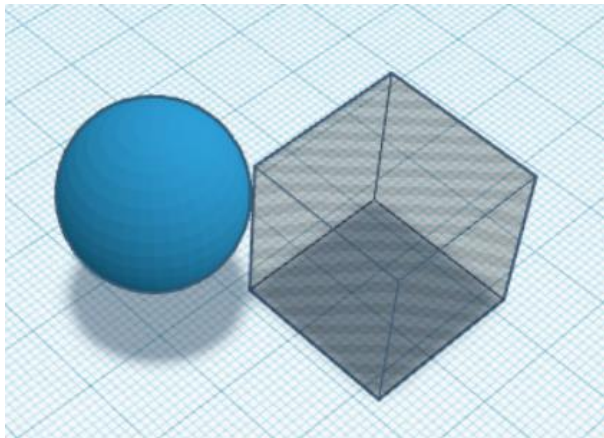


Check if the ball is contacting the face of the octant

```
if np.sum((query_offset_abs < octant.extent).astype(np.int)) >= 2:  
    return True
```



Function *overlaps* – Case 3.1



```
# consider the case that the ball is contacting the edge or corner of the octant
# since the case of the ball center (query) inside octant has been considered,
# we only consider the ball center (query) outside octant
x_diff = max(query_offset_abs[0] - octant.extent, 0)
y_diff = max(query_offset_abs[1] - octant.extent, 0)
z_diff = max(query_offset_abs[2] - octant.extent, 0)

return x_diff * x_diff + y_diff * y_diff + z_diff * z_diff < radius * radius
```

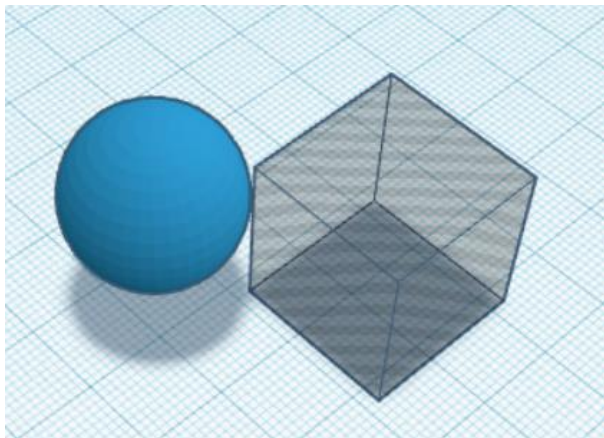



Function *overlaps* – Case 3.2

In 3D, there is the case that the cube' s edge cut into the query ball

```
# consider the case that the ball is contacting the edge or corner of the octant
# since the case of the ball center (query) inside octant has been considered,
# we only consider the ball center (query) outside octant
x_diff = max(query_offset_abs[0] - octant.extent, 0)
y_diff = max(query_offset_abs[1] - octant.extent, 0)
z_diff = max(query_offset_abs[2] - octant.extent, 0)

return x_diff * x_diff + y_diff * y_diff + z_diff * z_diff < radius * radius
```



That's why there is a “*max*” to reduce this case into 3.1



Octree Radius NN Search



Simple one: replace KNNResultSet with RadiusNNResult Set

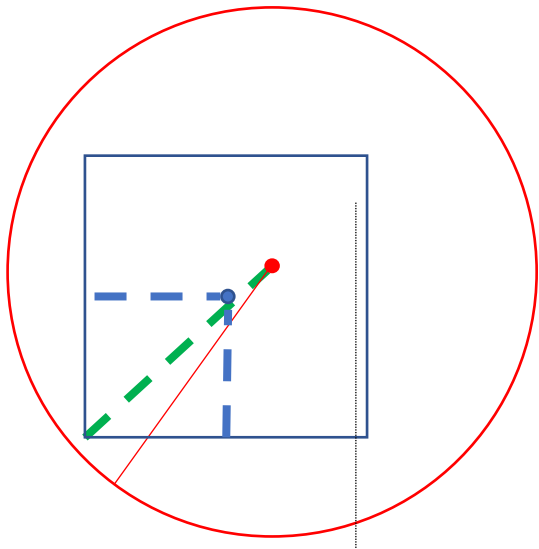


Better one:

- If the query ball *contains* the octant, just compare the query with all point
- No need to go into children of that octant



Function *contains*



```
def contains(query: np.ndarray, radius: float, octant: Octant):  
    """  
    Determine if the query ball contains the octant  
    :param query:  
    :param radius:  
    :param octant:  
    :return:  
    """  
  
    query_offset = query - octant.center  
    query_offset_abs = np.fabs(query_offset)  
  
    query_offset_to_farthest_corner = query_offset_abs + octant.extent  
    return np.linalg.norm(query_offset_to_farthest_corner) < radius
```

Green dash line

Red line



Octree Search Complexity



1NN search is $O(\log n)$



kNN/radiusNN complexity is hard to analyze

- Depends on the distribution of points
- Depends on k or r
- Varies from $O(\log n)$ to $O(n)$

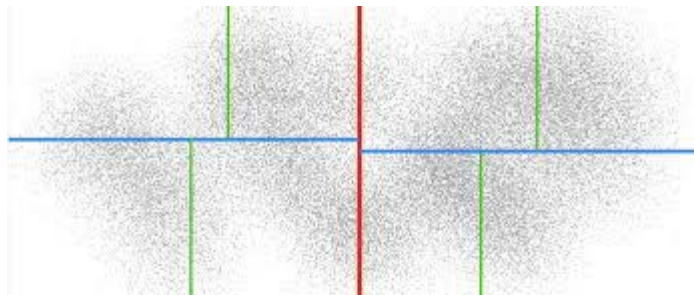
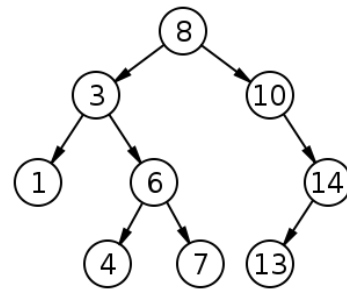


BST / Kd-tree / Octree



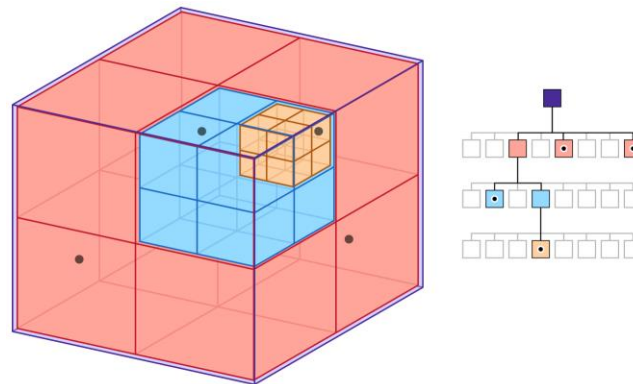
Dimension

- BST for one dimension
- Kd-tree works for any dimension
- Octree is optimized for 3D



Idea

- Same – space partition





Summary



Space partition



Find a method to skip some partitions



Pythons codes:

<https://github.com/lijx10/NN-Trees>



Homework

- We provide one $N \times 3$ point cloud
- 8-NN search for each point to the point cloud
- Implement 3 NN algorithms
 1. Numpy brute-force search
 2. `scipy.spatial.KDTree`
 3. Your own kd-tree/octree in python or C++
- Report timing using method 1 as baseline
- This is a competition!
 - Timing of method 3 determine your grade

感谢聆听 !

Thanks for Listening